

Parallel Performance of Molecular Dynamics Trajectory Analysis

Mahzad Khoshlessan^a, Ioannis Paraskevatos^b, Geoffrey C. Fox^c, Shantenu Jha^b,
Oliver Beckstein^{a,d,*}

^a*Department of Physics, Arizona State University, Tempe, AZ 85281, USA*

^b*Department of Electrical & Computer Engineering, Rutgers University, Piscataway, NJ 08854, USA*

^c*Digital Science Center, Indiana University, Bloomington, IN 47405*

^d*Center for Biological Physics, Arizona State University, Tempe, AZ 85281, USA*

Abstract

The performance of biomolecular molecular dynamics (MD) simulations has steadily increased on modern high performance computing (HPC) resources but acceleration of the analysis of the output trajectories has lagged behind so that analyzing simulations is increasingly becoming a bottleneck. To close this gap, we studied the performance of parallel trajectory analysis with MPI and the Python *MDAnalysis* library on three different XSEDE supercomputers where trajectories were read from a Lustre parallel file system. We found that strong scaling performance was impeded by stragglers, MPI processes that were slower than the typical process and that therefore dominated the overall run time. Stragglers were less prevalent for compute-bound workloads, thus pointing to file reading as a crucial bottleneck for scaling. However, a more complicated picture emerged in which both the computation and the ingestion of data exhibited close to ideal strong scaling behavior whereas stragglers were primarily caused by either large MPI communication costs or long times to open the single shared trajectory file. We improved overall strong scaling performance by two different approaches to file access, namely subfilings (splitting the trajectory into as many trajectory segments as number of processes) and MPI-IO with Parallel HDF5 trajectory files. Applying these strategies, we obtained near ideal strong scaling on up to 384 cores (16 nodes).

*Corresponding author

Email addresses: mkhoshle@asu.edu (Mahzad Khoshlessan), i.paraskev@rutgers.edu (Ioannis Paraskevatos), gcf@indiana.edu (Geoffrey C. Fox), shantenu.jha@rutgers.edu (Shantenu Jha), oliver.beckstein@asu.edu (Oliver Beckstein)

We summarize our lessons-learned in guidelines and strategies on how to take advantage of the available HPC resources to gain good scalability and potentially reduce trajectory analysis times by two orders of magnitude compared to the prevalent serial approach.

Keywords: Python, MPI, HPC, MDAnalysis, MPI I/O, Global Arrays, HDF5, Straggler, Molecular Dynamics, Big Data, Trajectory Analysis

1. Introduction

Molecular dynamics (MD) simulations are a powerful method to generate new insights into the function of biomolecules [1–5]. These simulations produce trajectories—time series of atomic coordinates—that now routinely include millions of time steps and can measure Terabytes in size. These trajectories need to be analyzed using statistical mechanics approaches [6, 7] but because of the increasing size of data, trajectory analysis is becoming a bottleneck in typical biomolecular simulation scientific workflows [8]. Many data analysis tools and libraries have been developed to extract the desired information from the output trajectories from MD simulations [9–22] but few can efficiently use modern High Performance Computing (HPC) resources to accelerate the analysis stage. MD trajectory analysis primarily requires *reading* of data from the file system; the processed output data are typically negligible in size compared to the input data and therefore we exclusively investigate the reading aspects of trajectory I/O (i.e., the “T”). We focus on the *MDAnalysis* package [17, 18], which is an open-source object-oriented Python library for structural and temporal analysis of MD simulation trajectories and individual protein structures. Although *MDAnalysis* accelerates selected algorithms with OpenMP, it is not clear how to best use it for scaling up analysis on multi-node supercomputers. Here we discuss the challenges and lessons-learned for making parallel analysis on HPC resources feasible with *MDAnalysis*, which should also be broadly applicable to other general purpose trajectory analysis libraries.

Previously, we had used a parallel split-apply-combine approach to study the performance of the commonly performed “RMSD fitting” analysis problem [23, 24],

24 which calculates the minimal root mean squared distance (RMSD) of the positions
25 of a subset of atoms to a reference conformation under optimization of rigid body
26 translations and rotations [7, 25, 26]. We had investigated two parallel implementa-
27 tions, one using *Dask* [27] and one using the message passing interface (MPI) with
28 *mpi4py* [28, 29]. For both *Dask* and MPI, we had previously only been able to obtain
29 good strong scaling performance within a single node. Beyond a single node perfor-
30 mance had dropped due to *straggler* tasks, a subset of tasks that had performed abnor-
31 mally slower than the typical task execution times; the total execution time had become
32 dominated by stragglers and overall performance had decreased. Stragglers are a well-
33 known challenge to improving performance on HPC resources [30] but there has been
34 little discussion of their impact in the biomolecular simulation community.

35 In the present study, we analyzed the MPI case in more detail to better understand
36 the origin of stragglers with the goal to find parallelization approaches to speed up
37 parallel post-processing of MD trajectories in the *MDAnalysis* library. We especially
38 wanted to make efficient use of the resources provided by current supercomputers such
39 as multiple nodes with hundreds of CPU cores and a Lustre parallel file system.

40 As in our previous study [23] we selected the commonly used RMSD algorithm
41 implemented in *MDAnalysis* as a typical use case. We employed the single program
42 multiple data (SPMD) paradigm to parallelize this algorithm on three different HPC
43 resources (XSEDE’s *SDSC Comet*, *LSU SuperMic*, and *PSC Bridges* [31]). With
44 SPMD, each process executes essentially the same operations on different parts of
45 the data. The three clusters differed in their architecture but all used Lustre as their
46 parallel file system. We used Python (<https://www.python.org/>), a machine-
47 independent, byte-code interpreted, object-oriented programming language, which is
48 well-established in the biomolecular simulation community with good support for par-
49 allel programming for HPC [28, 32]. We found that communication and reading I/O
50 were the two main scalability bottlenecks, with some indication that read I/O might
51 have been interfering with the communications. We therefore focused on two different
52 approaches to mitigate I/O bottlenecks: MPI parallel I/O (MPI-IO) with the HDF5 file
53 format and subfiling (trajectory file splitting). For subfiling, we obtained good results
54 with the *Global Arrays* package [32, 33], which provides a convenient layer to access

55 and manage arrays over multiple MPI ranks. Both MPI-IO and subfiling eliminated
56 stragglers and improved the performance with near ideal scaling, $S(N) = N$, i.e., the
57 speed-up S scaled linearly with the number N of CPU cores while exhibiting a slope
58 of one.

59 The paper is organized as follows: We first review stragglers and existing ap-
60 proaches to parallelizing MD trajectory analysis in section 2. We describe the software
61 packages and algorithms in section 3 and the benchmarking environment in section
62 4. Section 5 explains how we measured performance. The main results are presented
63 in section 6, with section 7 demonstrating reproducibility on different supercomput-
64 ers. We provide general guidelines and lessons-learned in section 8 and finish with
65 conclusions in section 9.

66 2. Background and Related Work

67 In our previous work, we found that straightforward implementation of simple par-
68 allelization with a split-apply-combine algorithm in Python failed to scale beyond a
69 single compute node [23] because a few tasks (MPI-ranks or Dask [27] processes) took
70 much longer than the typical task and so limited the overall performance. However, the
71 cause for these *straggler* tasks remained obscure. Here, we studied the straggler prob-
72 lem in the context of an MPI-parallelized trajectory analysis algorithm in Python and
73 investigated solutions to overcome it. We briefly review stragglers in section 2.1 and
74 summarize existing approaches to parallel trajectory analysis in section 2.2.

75 2.1. Stragglers

76 *Stragglers* or *outliers* were traditionally considered in the context of MapReduce
77 jobs that consist of multiple tasks that all have to finish for the job to succeed: A
78 straggler was a task that took an “unusually long time to complete” [34] and there-
79 fore substantially impeded job completion. In general, any component of a parallel
80 workflow whose runtime exceeds a typical run time (for example, 1.5 times the median
81 runtime) can be considered a straggler [35]. Stragglers are a challenge for improving
82 performance on HPC resources [30]; they are a known problem in frameworks such

83 as MapReduce [34, 35], Spark [36–39], Hadoop [34], cloud data centers [30, 40], and
 84 have a high impact on performance and energy consumption of big data systems [41].
 85 Both internal and external factors are known to contribute to stragglers. Internal fac-
 86 tors include heterogeneous capacity of worker nodes and resource competition due to
 87 other tasks running on the same worker node. External factors include resource com-
 88 petition due to co-hosted applications, input data skew, remote input or output source
 89 being too slow, faulty hardware [34, 42], and node mis-configuration [34]. Competition
 90 over scarce resources [35], in particular the network bandwidth, was found to lead to
 91 stragglers in writing on Lustre file systems [43]. Garbage collection [36, 37], Java vir-
 92 tual machine (JVM) positioning to cores [36], delays introduced while the tasks move
 93 from the scheduler to execution [38], disk I/O during shuffling, Java’s just-in-time com-
 94 pilation [37], output skew [37], high CPU utilization, disk utilization, unhandled I/O
 95 access requests, and network package loss [30] were also among other external factors
 96 that might introduce stragglers. A wide variety of approaches have been investigated
 97 for detecting and mitigating stragglers, including tuning resource allocation and par-
 98 allelism such as breaking the workload into many small tasks that are dynamically
 99 scheduled at runtime [44], slow Node-Threshold [34], speculative execution [34] and
 100 cause/resource-aware task management [35], sampling or data distribution estimation
 101 techniques, SkewTune to avoid data imbalance [45], dynamic work rebalancing [40],
 102 blocked time analysis [46], and intelligent scheduling [47].

103 In the present study, we analyzed large MD trajectories in parallel with MPI and
 104 Python and observed large variations in the completion time of individual MPI ranks.
 105 These variations bore some similarity to the straggler tasks observed in MapReduce
 106 frameworks so we approached analyzing and eliminating them in a similar fashion
 107 by systematically looking at different components of the problem, including read I/O
 108 from the shared Lustre file system and MPI communication. Even though we specifi-
 109 cally worked in with the *MDAnalysis* package, all these principles and techniques are
 110 potentially applicable to MPI-parallelized data analysis in other Python-based libraries.

111 2.2. Other Packages with Parallel Analysis Capabilities

112 Different approaches to parallelizing the analysis of MD trajectories have been pro-
113 posed. HiMach [14] introduces scalable and flexible parallel Python framework to deal
114 with massive MD trajectories, by combining and extending Google’s MapReduce and
115 the VMD analysis tool [11]. HiMach’s runtime is responsible to parallelize and dis-
116 tribute Map and Reduce classes to assigned cores. HiMach uses parallel I/O for file
117 access during map tasks and MPI_Allgather in the reduction process. HiMach, how-
118 ever, does not discuss parallel analysis of analysis types that cannot be implemented
119 via MapReduce. Furthermore, HiMach is not available under an open source license,
120 which makes it difficult to integrate community contributions and add new state-of-the-
121 art methods.

122 Wu et. al. [48] present a scalable parallel framework for distributed-memory post-
123 simulation data analysis. This work consists of an interface that allows a user to write
124 analysis programs sequentially, and the machinery that ensures these programs exe-
125 cute in parallel automatically. The main components of the proposed framework are
126 (1) domain decomposition that splits computational domain into blocks with specified
127 boundary conditions, (2) HDF5 based parallel I/O (3) data exchange that communi-
128 cates ghost atoms between neighbor blocks, and (4) parallel analysis implementation
129 of a real-world analysis application. This work does not discuss analysis methods
130 which cannot be implemented using MapReduce and is limited to HDF5 file format.

131 Zazen [49] is a novel task-assignment protocol to overcome the I/O bottleneck for
132 many I/O bound tasks. This protocol caches a copy of simulation output files on the
133 local disks of the compute nodes of a cluster, and uses co-located data access with com-
134 putation. Zazen is implemented in a parallel disk cache system and avoids the overhead
135 associated with querying metadata servers by reading data in parallel from local disks.
136 This approach has also been used to improve the performance of HiMach [14]. It, how-
137 ever, advocates a specific architecture where a parallel supercomputer, which runs the
138 simulations, immediately pushes the trajectory data to a local analysis cluster where
139 trajectory fragments are cached on node-local disks. In the absence of such a specific
140 workflow, one would need to stage the trajectory across nodes, and the time for data
141 distribution is likely to reduce any gains from the parallel analysis.

142 VMD [11, 50] provides molecular visualization and analysis tool through algorithmic and memory efficiency improvements, vectorization of key CPU algorithms, GPU
143 analysis and visualization algorithms, and good parallel I/O performance on supercom-
144 puters. It is one of the most advanced programs for the visualization and analysis of
145 MD simulations. It is, however, a large monolithic program, that can only be driven
146 through its built-in Tcl interface and thus is less well suited as a library that allows the
147 rapid development of new algorithms or integration into workflows.

149 CPPTraj [19] offers multiple levels of parallelization (MPI and OpenMP) in a
150 monolithic C++ implementation. CCPTraj allows parallel reads between frames of
151 the same trajectory but is especially geared towards processing an ensemble of many
152 trajectories in parallel.

153 pyPcazip [51] is a suite of software tools written in Python for compression and
154 analysis of MD simulation data, in particular ensembles of trajectories. pyPcazip is
155 MPI parallelised and is specific to PCA-based investigations of MD trajectories and
156 supports a wide variety of trajectory file formats (based on the capabilities of the un-
157 derlying mdtraj package [20]). pyPcazip can take one or many input MD trajectory files
158 and convert them into a highly compressed, HDF5-based pcz format with insignificant
159 loss of information. However, the package does not support general purpose analysis.

160 *In situ* analysis is an approach to execute analysis simultaneously with the running
161 MD simulation so that I/O bottlenecks are mitigated [52, 53]. Malakar *et al.* studied
162 the scalability challenges of time and space shared modes of analyzing large-scale MD
163 simulations through a topology-aware mapping for simulation and analysis using the
164 LAMMPS code [52]. Similarly, Taufer and colleagues [53] presented their own frame-
165 work for *in situ* analysis, which is based on the fast on-the-fly calculation of metadata
166 that characterizes protein substructures via maximum eigenvalues of distance matri-
167 ces. These metadata are used to index trajectory frames and enable targeted analysis of
168 trajectory subsets. Both studies provide important ideas and approaches towards mov-
169 ing towards online-analysis in conjunction with a running simulation but are limited in
170 generality.

171 All of the above frameworks provide tools for parallel analysis of MD trajectories.
172 These frameworks, however, tend to fall short in providing parallelism in the context

173 of a general and flexible library for the analysis of MD trajectories. Although straggler
174 tasks are a common challenge arising in parallel analysis and are well-known in the
175 data analysis community (see Section 2.1), there is, to our knowledge, little discussion
176 about this problem in the biomolecular simulation community. Our own experience
177 with a MapReduce approach in *MDAnalysis* [23] suggested that stragglers might be a
178 somewhat under-appreciated problem. Therefore, in the present work we want to better
179 understand requirements for efficient parallel analysis of MD trajectories in *MDAnaly-*
180 *sis*, but to also provide more general guidance that could benefit developments in other
181 libraries inside and outside of the scope of analysis of MD simulations.

182 3. Algorithms and Software Packages

183 For our investigation of parallel trajectory analysis we focus on using MPI as the
184 standard approach to parallelization in HPC. We employ the Python language, which
185 is widely used in the scientific community because it facilitates rapid development of
186 small scripts and code prototypes as well as development of large applications and
187 highly portable and reusable modules and libraries. We use the *MDAnalysis* library to
188 calculate a “RMSD timeseries” (explained in section 3.1) as a representative use case.
189 Further details on the software packages are provided in sections 3.2–3.4.

190 3.1. RMSD Calculation with *MDAnalysis*

191 Simulation data exist in trajectories in the form of time series of atom positions and
192 sometimes velocities. Trajectories come in a plethora of different and idiosyncratic file
193 formats. *MDAnalysis* [17, 18] is a widely used open source library to analyze trajectory
194 files with an object oriented interface. The library is written in Python, with time
195 critical code in C/C++/Cython. *MDAnalysis* supports most file formats of simulation
196 packages including CHARMM [54], Gromacs [55], Amber [56], and NAMD [57] and
197 the Protein Data Bank [58] format. At its core, it reads trajectory data in different
198 formats and makes them available through a uniform API; specifically, coordinates are
199 represented as standard NumPy arrays [59].

As a test case that is representative of a common task in the analysis of biomolecular
simulation trajectories we calculated the timeseries of the minimal structural root mean

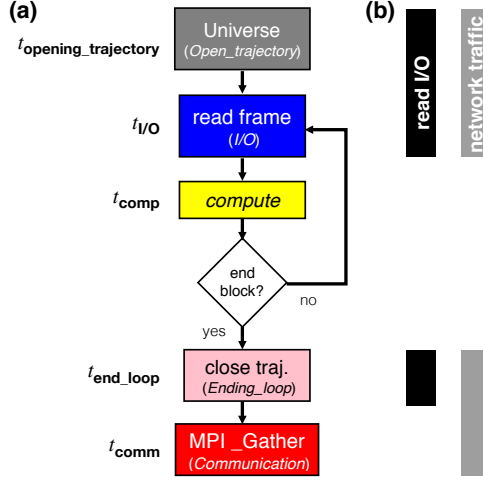


Figure 1: Flow chart of the MPI-parallelized RMSD algorithm, Algorithm 1. (a) Each MPI process performs the same steps but reads trajectory frames from different blocks of the trajectory. The color scheme and labels in italics correspond to the colors and labels for measured timing quantities in the following graphs (e.g., Figs. 2c and 2d). The names of the corresponding timing quantities from Table 3 are listed next to each step. (b) Steps that access the shared Lustre file system with read I/O are included in the black bars; steps that communicate via the shared InfiniBand network are included in the gray bars. The Lustre file system is accessed through the network and hence all I/O steps also use the network.

square distance (**RMSD**) after rigid body superposition [7, 26]. The RMSD is used to show the rigidity of protein domains and more generally characterizes structural changes. It is calculated as a function of time t as

$$\text{RMSD}(t) = \min_{\mathbf{R}, \mathbf{t}} \sqrt{\frac{1}{N} \sum_{i=1}^N \left[(\mathbf{R} \cdot \mathbf{x}_i(t) + \mathbf{t}) - \mathbf{x}_i^{\text{ref}} \right]^2} \quad (1)$$

200 where $\mathbf{x}_i(t)$ is the position of atom i at time t , $\mathbf{x}_i^{\text{ref}}$ is its position in a reference structure
 201 and the distance between these two is minimized by finding the optimum 3×3 rotation
 202 matrix \mathbf{R} and translation vector \mathbf{t} . The optimum rigid body superposition was calculated
 203 with the QCPROT algorithm [25, 60] (implemented in Cython and available through
 204 the `MDAnalysis.analysis.rms` module [18]).

205 The RMSD trajectory analysis was parallelized as outlined in the flow chart in
 206 Figure 1, with further details available in Algorithm 1. Each MPI process loads the
 207 core MDAnalysis structure (called the `Universe`), which includes loading a shared
 208 “topology” file with the simulation system information and opening the shared trajec-

209 tory file. Each process operates on a different block of frames and iterates through them
 210 by reading the coordinates of a single frame into memory and performing the RMSD
 211 computation with them. Once all frames in the block are processed, the trajectory file
 212 is closed and results are communicated to MPI rank 0 using `MPI_Gather()`.

213 The RMSD was determined for a subset of protein atoms, the $N = 214$ C_α atoms
 214 of our test system (see section 4.3 for details). The time complexity for the RMSD
 215 Algorithm 1 is $O(T \times N^2)$ [25] where T is the number of frames in the trajectory and
 216 N the number of particles included in the RMSD calculation.

Algorithm 1 MPI-parallel Multi-frame RMSD Algorithm

Input: *size*: Total number of frames
ref: mobile group in the initial frame which will be considered as reference
start & stop: Starting and stopping frame index
topology & trajectory: files to read the data structure from
Output: Calculated RMSD arrays

```

1: procedure Block_RMSD(topology, trajectory, ref, start, stop)
2:    $u \leftarrow \text{Universe}(\text{topology}, \text{trajectory})$  ▷  $u$  hold all the information of the physical system
3:    $g \leftarrow u.\text{frames}[\text{start}:\text{stop}]$ 
4:   for  $\forall i\text{frame}$  in  $g$  do
5:      $\text{results}[i\text{frame}] \leftarrow \text{RMSD}(g, \text{ref})$ 
6:   end for
7:   return results
8: end procedure
9:
10: MPI Init
11: rank  $\leftarrow$  rank ID
12: index  $\leftarrow$  indices of mobile atom group
13: xref0  $\leftarrow$  Reference atom group's position
14: out  $\leftarrow$  Block_RMSD(topology, trajectory, xref0, start=start, stop=stop)
15:
16: Gather(out, RMSD_data, rank_ID=0)
17: MPI Finalize

```

217 **3.2. MPI for Python (*mpi4py*)**

218 MPI for Python (*mpi4py*) is a Python wrapper for the Message Passing Interface
 219 (MPI) standard and allows any Python program to employ multiple processors [28,
 220 29]. Performance degradation due to using *mpi4py* is not prohibitive [28, 29] and
 221 the overhead is far smaller than the overhead associated with the use of interpreted
 222 versus compiled languages [32]. Overheads in *mpi4py* are small compared to C code
 223 if efficient raw memory buffers are used for communication [28], as used in the present
 224 study.

225 3.3. *Global Arrays Toolkit*

226 The *Global Arrays* (GA) toolkit provides users with a language interface that al-
227 lows them to distribute data while maintaining the type of global index space and pro-
228 gramming syntax similar to what is available when programming on a single proces-
229 sor [33]. *Global Arrays* is implemented with Fortran-77 and C bindings and provides
230 C++ and Python interfaces. It allows manipulating physically distributed dense multi-
231 dimensional arrays without explicitly defining communication and synchronization be-
232 tween processes. The underlying communication is determined by a runtime environ-
233 ment, which defaults to the *Communication runtime for Extreme Scale* (ComEx) [61].
234 ComEx uses shared memory for intra-node communication and inter-node communi-
235 cation employs ComEx with MPI. *Global Arrays in NumPy* (GAIN) extends GA to
236 Python through Numpy [32]. The *Global Arrays* toolkit provides functions to create
237 global arrays (`ga_create()`) and to copy data to (`ga_put()`) and from (`ga_get()`)
238 such a global array, as well as additional functions for copying between arrays and free-
239 ing them [32]. When a global array is created (`ga_create()`) each process will cre-
240 ate an array of the same shape and size, physically located in the local memory space
241 of that process [33]. The GA library maintains a list of all these memory locations,
242 which can be queried with the `ga_access()` function. Using a pointer returned by
243 `ga_access()`, one can directly modify the data that is local to each process. When
244 a process tries to access a block of data the request is first decomposed into individ-
245 ual blocks representing the contribution to the total request from the data held locally
246 on each process (*B. J. Palmer and J. Daily, personal communication*). The requesting
247 process then makes individual requests to each of the remote processes.

248 GA allows independent, asynchronous, and efficient access to logical blocks of
249 physically distributed arrays, with no need for explicit cooperation by other processes;
250 in particular, it allows data locality to be explicitly specified and used [62]. We inves-
251 tigated if communication cost could be reduced by using *Global Arrays*. Algorithm 2
252 describes the RMSD algorithm with *Global Arrays* instead of MPI.

Algorithm 2 MPI-parallel Multi-frame RMSD using Global Arrays

Input: *size*: Total number of frames assigned to each rank N_b
g_a: Initialized Global Arrays
xref0: mobile group in the initial frame which will be considered as reference
start & *stop*: that tell which block of trajectory (frames) is assigned to each rank
topology & *trajectory*: files to read the data structure from
Include: `Block_RMSD()` from Algorithm 1

```
1: bsize  $\leftarrow$  ceil(trajectory.number_frames / size)
2: g_a  $\leftarrow$  ga.create(ga.C_DBL, [bsize*size,2], "RMSD")
3: buf  $\leftarrow$  np.zeros([bsize*size,2], dtype=float)
4: out  $\leftarrow$  Block_RMSD(topology, trajectory, xref0, start=start, stop=stop)
5: ga.put(g_a, out, (start,0), (stop,2))
6: if rank == 0 then
7:   buf  $\leftarrow$  ga.get(g_a, lo=None, hi=None)
8: end if
```

3.4. MPI and Parallel HDF5

HDF5 is a structured self-describing hierarchical data format which is the standard mechanism for storing large quantities of numerical data in Python (<http://www.hdfgroup.org/HDF5>, [63]). Parallel HDF5 (*PHDF5*) typically sits on top of a MPI-IO layer and can use MPI-IO optimizations. In *PHDF5*, all file access is coordinated by the MPI library; otherwise, multiple processes would compete over accessing the same file on disk. MPI-based applications launch multiple parallel instances of the Python interpreter that communicate with each other via the MPI library. Implementation is straightforward as long as the user supplies a MPI communicator and takes into account some constraints required for data consistency [63]. *HDF5* itself handles nearly all the details involved with coordinating file access when the shared file is opened through the *mpio* driver.

MPI has two flavors of operation: collective (all processes have to participate in the same order) and independent (processes can perform the operation in any order or not at all) [63]. With *PHDF5*, modifications to file metadata must be performed collectively and although all processes perform the same task, they do not need to be synchronized [63]. Other tasks and any type of data operations can be performed independently by processes. In the present study, we use independent operations.

4. Benchmark Environment

Our benchmark environment consisted of three different XSEDE [31] HPC resources (described in section 4.1), the software stack used (section 4.2), which had to be compiled for each resource, and the common test data set (section 4.3).

4.1. HPC Resources

The computational experiments were executed on standard compute nodes of three XSEDE [31] supercomputers, *SDSC Comet*, *PSC Bridges*, and *LSU SuperMIC* (Table 1). *SDSC Comet* is a 2 PFlop/s cluster with 2,020 compute nodes in total. It is optimized for running a large number of medium-size calculations (up to 1,024 cores) to support the most prevalent type of calculation on XSEDE resources. *PSC Bridges* is a 1.35 PFlop/s cluster with different types of computational nodes, including 16 GPU nodes, 8 large memory and 2 extreme memory nodes, and 752 regular nodes. It was designed to flexibly support both traditional (medium scale calculations) and non-traditional (data analytics) HPC uses. *LSU SuperMIC* offers 360 standard compute nodes with a peak performance of 557 TFlop/s. The parallel file system on all three machines is Lustre (<http://lustre.org/>) and is shared between the nodes of each cluster.

Name	Nodes	Number of Nodes	CPUs	RAM	Network Topology	Scheduler and Resource Manager	parallel file system
<i>SDSC Comet</i>	Compute	6400	2 Intel Xeon (E5-2680v3) 12 cores/CPU, 2.5 GHz	128 GB DDR4 DRAM	56 Gbps IB	SLURM	Lustre
<i>PSC Bridges</i>	RSM	752	2 Intel Haswell (E5-2695 v3) 14 cores/CPU, 2.3 GHz	128 GB, DDR4-2133Mhz	12.37 Gbps OPA	SLURM	Lustre
<i>LSU SuperMIC</i>	Standard	360	2 Intel Ivy Bridge (E5-2680) 10 cores/CPU, 2.8 GHz	64 GB, DDR3-1866Mhz	56 Gbps IB	PBS	Lustre

Table 1: Configuration of the HPC resources that were benchmarked. Only a subset of the total available nodes were used. IB: InfiniBand; OPA: Omni-Path Architecture.

4.2. Software

Table 2 lists the tools and libraries that were required for our computational experiments. Many domain specific packages are not available in the standard software installation on supercomputers. We therefore had to compile them, which in some cases required substantial effort due to non-standard building and installation procedures or lack of good documentation. Because this is a common problem that hinders reproducibility we provide detailed version information, notes on the installation process, as well as comments on the ease of installation and the quality of the documentation in Table 2. For the MPI implementation we used Open MPI release 1.10.7 (<https://www.open-mpi.org/>) consistently everywhere. Detailed instructions to create the computing environments together with the benchmarking code can be

found in the GitHub repository. Carefully setting up the same software stack on the three different supercomputers allowed us to clearly demonstrate the reproducibility of our results and showed that our findings were not dependent on machine specifics.

Package	Version	Description	Ease of Installation	Documentation	Installation	Dependencies
GCC	4.9.4	GNU Compiler Collection	0	++	via configuration files, environment or command line options, minimal configuration is required	–
Open MPI	1.10.7	MPI Implementation	0	++	via configuration files, environment or command line options, minimal configuration is required	–
Global Arrays	5.6.1	Global Arrays	–	+	via configuration files, environment or command line options, several optional configuration settings available	MAMA, ARMCI MPI 1.x/2.x/3.x implementation like Open MPI built with shared/dynamic libraries, GCC
Python	2.7.13	Python language	+	++	Conda Installation	–
MPI4py	3.0.0	MPI for Python	+	++	Conda Installation	Python 2.7 or above, MPI 1.x/2.x/3.x implementation like Open MPI built with shared/dynamic libraries, Cython
GA4py	1.0	Global Arrays for Python	0	0	Python Setuptools	Global Arrays, Python 2 only, MPI 1.x/2.x/3.x implementation like Open MPI built with shared/dynamic libraries, Cython, MPI4py, Numpy
PHDF5	1.10.1	Parallel HDF5	–	++	via configuration files, environment or command line options, several optional configuration settings available	MPI 1.x/2.x/3.x implementation like Open MPI GNU, MPIF90, MPICC, MPICXX
H5py	2.7.1	Pythonic wrapper around the HDF5	+	++	Conda Installation	Python 2.7, or above, PHDF5, Cython
MDAnalysis	0.17.0	Python library to analyze trajectories from MD simulations	+	++	Conda Installation	Python >=2.7, Cython, GNU, Numpy

Table 2: Detailed comparison on the dependencies and installation of different software packages used in the present study. Software was built from source or obtained via a package manager and installed on the multi-user HPC systems in Table 1. Evaluation of ease of installation and documentation uses a subjective scale with “++” (excellent), “+” (good), “0” (average), and “–” (difficult/lacking) and reflects the experience of a typical domain scientist at the graduate/post-graduate level in a discipline such as computational biophysics or chemistry.

4.3. Data Set

The test system contained the protein adenylate kinase with 214 amino acid residues and 3341 atoms in total [64] and the topology information (atoms types and bonds) was stored in a file in CHARMM PSF format. The test trajectory was created by concatenating 600 copies of a MD trajectory with 4,187 time frames (saved every 240 ps for a total simulated time of 1.004 μ s) in CHARMM DCD format [65] and converting to Gromacs XTC format trajectory, as described for the “600x” trajectory

in Khoshlessan et al. [23]. The trajectory had a file size of about 30 GB and contained 2,512,200 frames (corresponding to 602.4 μ s simulated time). The file size was relatively small because water molecules that were also part of the original MD simulations were stripped to reduce the original file size by a factor of about 10; such preprocessing is a common approach if one is only interested in the protein behavior. Thus, the trajectory represents a small to medium system size in the number of atoms and coordinates that have to be loaded into memory for each time frame. The XTC format is a format with lossy compression [66, 67], which also contributed to the compact file size. XTC trades lower I/O demands for higher CPU demands during decompression and therefore performed well in our previous study [23]. Although 2,512,200 frames represents a long simulation for current standards, such trajectories will become increasingly common due to the use of special hardware [68, 69] and GPU-acceleration [55, 70, 71].

5. Methods

Documentation and benchmark codes are made available in the code repository <https://github.com/hpcanalytics/supplement-hpc-py-parallel-mdanalysis> under the GNU General Public License v3.0 (code) and the Creative Commons Attribution-ShareAlike (documentation). These materials should enable users to recreate the computational environment on the tested XSEDE HPC resources (*SDSC Comet*, *PSC Bridges*, *LSU SuperMIC*), prepare data files, and run the computational experiments.

In the following we define the quantities and approach used for our performance measurements, with a full summary of all definitions in Table 3. We evaluated MPI performance of the parallel RMSD timeseries algorithm 1 by timing the total time to solution as well as the execution time for different parts of the code for individual MPI ranks with the help of the Python `time.time()` function.

5.1. Timing Observables

We abbreviate the timings in the following as variables t_{Ln} where Ln refers to the line number in algorithm 1. We measured in the function `block_rmsd()` the *read I/O*

Quantity	Definition
N_b	$N_{\text{frames}}^{\text{total}}/N$
$t_{\text{end_loop}}$	t_{L6}
$t_{\text{opening_trajectory}}$	$t_{L2} + t_{L3}$
t_{comp}	$\sum_{\text{frame}=1}^{N_b} t_{\text{comp}}^{\text{frame}}$
$t_{I/O}$	$\sum_{\text{frame}=1}^{N_b} t_{I/O}^{\text{frame}}$
$t_{\text{all_frame}}$	$t_{L4} + t_{L5} + t_{L6}$
t_{RMSD}	$t_{L1} + \dots + t_{L8}$
$t_{\text{comm/MPI}}$	t_{L16}
$t_{\text{comm/GA}}$	$t_{L5} + t_{L6} + t_{L7} + t_{L8}$
t_{comm}	$t_{\text{comm/MPI}}$ (Alg. 1) or $t_{\text{comm/GA}}$ (Alg. 2)
$t_{\text{Overhead1}}$	$t_{\text{all_frame}} - t_{I/O} - t_{\text{comp}} - t_{\text{end_loop}}$
$t_{\text{Overhead2}}$	$t_{\text{RMSD}} - t_{\text{all_frame}} - t_{\text{opening_trajectory}}$
t_N	$t_{\text{RMSD}} + t_{\text{comm}}$
$\overline{t_{\text{comp}}}$	$\frac{1}{N} \sum_{\text{rank}=1}^N t_{\text{comp}}$
$\overline{t_{I/O}}$	$\frac{1}{N} \sum_{\text{rank}=1}^N t_{I/O}$
$\overline{t_{\text{comm}}}$	$\frac{1}{N} \sum_{\text{rank}=1}^N t_{\text{comm}}$
t_{total}	$\max t_N$

Table 3: Summary of measured timing quantities. Timings are collected for the specified line numbers in the code, labelled as t_{L_n} where L_n refers to the line number in the corresponding algorithm. $t_{\text{comm/MPI}}$ (in Algorithm 1) and $t_{\text{comm/GA}}$ (in Algorithm 2) are both referred to as t_{comm} in the text. Variables in the top half of the table refer to measurements of an individual MPI rank. Variables in the bottom half are aggregates such as averages over all ranks or the total time to solution.

337 *time* for ingesting the data of one trajectory frame from the file system into memory,
 338 $t_{I/O}^{\text{frame}} = t_{L4}$, and the *compute time* per trajectory frame to perform the computation,
 339 $t_{\text{comp}}^{\text{frame}} = t_{L5}$. The *total read I/O time for a MPI rank*, $t_{I/O} = \sum_{\text{frame}=1}^{N_b} t_{I/O}^{\text{frame}}$, is the sum
 340 over all I/O times for all the N_{frames} frames assigned to the rank; similarly, the *total*
 341 *compute time for a MPI rank* is $t_{\text{comp}} = \sum_{\text{frame}=1}^{N_b} t_{\text{comp}}^{\text{frame}}$. The time delay between the end
 342 of the last iteration and exiting the `for` loop is $t_{\text{end_loop}} = t_{L6}$. The time $t_{\text{opening_trajectory}} =$
 343 $t_{L2} + t_{L3}$ measures the problem setup, which includes data structure initialization and
 344 opening of topology and trajectory files. The *communication time*, $t_{\text{comm}} = t_{L16}$, is the
 345 time to gather all data from all processor ranks to rank zero. The total time (for all
 346 frames) spent in `block_rmsd()` is $t_{\text{RMSD}} = \sum_{i=1}^8 t_{L_i}$. There are parts of the code in
 347 `block_rmsd()` that are not covered by the detailed timing information of t_{comp} and
 348 $t_{I/O}$. Unaccounted time is considered as *overhead*. We define $t_{\text{Overhead1}}$ and $t_{\text{Overhead2}}$ as
 349 the overheads of the calculations (see Table 3 for the definitions); both are expected
 350 to be negligible, which was the case in all our measurements. Finally, the *total time*

351 to completion of a single MPI rank, when utilizing N cores for the execution of the
 352 overall experiment, is t_N , and as a result $t_{\text{RMSD}} + t_{\text{comm}} \equiv t_N$.

353 5.2. Performance Parameters

We measured the *total time to solution* $t_{\text{total}}(N)$ with N MPI processes on N cores, which is effectively $t_{\text{total}}(N) \approx \max(t_N)$. Strong scaling was quantified by the speed-up

$$S(N) = \frac{t_{\text{total}}(1)}{t_{\text{total}}(N)}, \quad (2)$$

relative to performance on a single core ($t_{\text{total}}(1)$), and the efficiency

$$E(N) = \frac{S(N)}{N}. \quad (3)$$

Averages over ranks were calculated as

$$\overline{t_{\text{comp}}} = \frac{1}{N} \sum_{\text{rank}=1}^N t_{\text{comp}} = \frac{1}{N} \sum_{\text{rank}=1}^N \sum_{\text{frame}=1}^{N_b} t_{\text{comp}}^{\text{frame}}, \quad (4)$$

$$\overline{t_{\text{I/O}}} = \frac{1}{N} \sum_{\text{rank}=1}^N t_{\text{I/O}} = \frac{1}{N} \sum_{\text{rank}=1}^N \sum_{\text{frame}=1}^{N_b} t_{\text{I/O}}^{\text{frame}}, \quad (5)$$

and

$$\overline{t_{\text{comm}}} = \frac{1}{N} \sum_{\text{rank}=1}^N t_{\text{comm}}. \quad (6)$$

Additionally, we introduced two performance parameters that we found to be indicative of the occurrence of stragglers. We defined the ratio of compute time to read I/O time for the serial code as

$$R_{\text{comp/I/O}} = \frac{t_{\text{comp}}}{t_{\text{I/O}}} = \frac{t_{\text{comp}}/N_{\text{frames}}^{\text{total}}}{t_{\text{I/O}}/N_{\text{frames}}^{\text{total}}} = \frac{\overline{t_{\text{comp}}^{\text{frame}}}}{\overline{t_{\text{I/O}}^{\text{frame}}}} \quad (7)$$

where the last equality shows that the ratio can also be computed from the average times per frame, $\overline{t_{\text{comp}}^{\text{frame}}}$ and $\overline{t_{\text{I/O}}^{\text{frame}}}$. $R_{\text{comp/I/O}}$ was calculated with the serial versions of our algorithms (on a single CPU core) in order to characterize the computational problem

in the absence of parallelization. The ratio of compute to communication time was defined by the ratio of average total compute time to the average total communication time

$$R_{\text{comp/comm}} = \frac{\overline{t_{\text{comp}}}}{\overline{t_{\text{comm}}}}. \quad (8)$$

Because t_{comm} cannot be measured for a serial code, we estimated $R_{\text{comp/comm}}$ from the rank-averages (Eqs. 4 and 6) for a given number of MPI ranks.


6. Computational Experiments

We had previously measured the performance of the MPI-parallelized RMSD analysis task on two different HPC resources (*SDSC Comet* and *TACC Stampede*) and had found that it only scaled well up to a single node due to high variance in the runtime of the MPI ranks, similar to the straggler phenomenon observed in big-data analytics [23]. However, the ultimate cause for this high variance could not be ascertained. We therefore performed more measurements with more detailed timing information (see section 5) on *SDSC Comet* (described in this section) and two other supercomputers (summarized in section 7) in order to better understand the origin of the stragglers and find solutions to overcome them.


6.1. RMSD Benchmark

We measured strong scaling for the RMSD analysis task (Algorithm 1) with the 2,512,200 frame test trajectory (section 4.3) on 1 to 72 cores (one to three nodes) of *SDSC Comet* (Figures 2a and 2b). We observed poor strong scaling performance beyond a single node (24 cores), comparable to our previous results [23]. A more detailed analysis showed that the RMSD computation, and to a lesser degree the read I/O, considered on their own, scaled well beyond 50 cores (yellow and blue lines in Figure 2c). But communication (sending results back to MPI rank 0 with `MPI_Gather()`; red line in Figure 2c) and the initial file opening (loading the system information into the `MDAnalysis.Universe` data structure from a shared “topology” file and opening the shared trajectory file; gray line in Figure 2c) started to dominate beyond 50 cores.


377 Communication cost and initial time for opening the trajectory were distributed un-
 378 evenly across MPI ranks, as shown in Figure 2d. The ranks that took much longer to
 379 complete than the typical execution time of the other ranks were the stragglers that hurt
 380 performance.




(a) Scaling total (five repeats)



(b) Speed-up (five repeats)



(c) Scaling for different components (five repeats)



(d) Time comparison on different parts of the calculations per MPI rank (example)

Figure 2: Performance of the RMSD task parallelized with MPI on *SDSC Comet*. Results were communicated back to rank 0. Five independent repeats were performed to collect statistics. (a-c) The error bars show standard deviation with respect to the mean. In serial, there is no communication and no data points are shown for $N = 1$ in (c). (d) Compute t_{comp} , read I/O $t_{\text{I/O}}$, communication t_{comm} , ending the for loop $t_{\text{end_loop}}$, opening the trajectory $t_{\text{opening_trajectory}}$, and overheads $t_{\text{overhead1}}$, $t_{\text{overhead2}}$ per MPI rank; see Table 3 for definitions. These are data from one run of the five repeats. MPI ranks 0, 12–27 and 29–72 are stragglers.

381 We qualitatively denoted by *straggler* any MPI rank that took at least about twice
382 as long as the group of ranks that finished fastest, roughly following the original de-
383 scription of a straggler as a task that took an “unusually long time to complete” [34].
384 The fast-finishing ranks were generally clearly distinguishable in the per-rank timings
385 such as in Figures 2d and A.11d. Such a qualitative definition of stragglers was suf-
386 ficient for our purpose to identify scalability bottlenecks, as shown in the following
387 discussion.

388 *Identification of Scalability Bottlenecks*

389 In the example shown in Figure 2d, 62 ranks out of 72 took about 60 s (the strag-
390 glers) whereas the remaining ranks only took about 20 s. In other instances, far fewer
391 ranks were stragglers, as shown, for example, in Figure A.11d. The detailed break-
392 down of the time spent on each rank (Figure 2d) showed that the computation, t_{comp} ,
393 was relatively constant across ranks. The time spent on reading data from the shared
394 trajectory file on the Lustre file system into memory, $t_{\text{I/O}}$, showed variability across dif-
395 ferent ranks. The stragglers, however, appeared to be defined by occasionally much
396 larger *communication* times, t_{comm} (line 16 in Algorithm 1), which were on the order of
397 30 s, and by larger times to initially open the trajectory (line 2 in Algorithm 1). t_{comm}
398 varied across different ranks and was barely measurable for a few of them. Although
399 the data in Figure 2d represented one run and in other instances different number of
400 ranks were stragglers, the averages over all ranks in five independent repeats (Fig-
401 ure 2c) showed that increased t_{comm} were generally the reason for large variations in
402 the run time for each rank. This initial analysis indicated that communication was a
403 major issue that prevented good scaling beyond a single node but the problems related
404 to file I/O also played an important role in limiting scaling performance.

405 *Influence of Hardware*

406 We ran the same benchmarks on multiple HPC systems that were equipped with
407 a Lustre parallel file system [XSEDE’s *PSC Bridges* (Fig. A.11) and *LSU SuperMIC*
408 (Fig. A.12)], and observed the occurrence of stragglers, in a manner very similar to the
409 results described for *SDSC Comet*. There was no clear pattern in which certain MPI

410 ranks would always be a straggler, and neither could we trace stragglers to specific
 411 cores or nodes. Therefore, the phenomenon of stragglers in the RMSD case was repro-
 412 ducible on different clusters and thus appeared to be independent from the underlying
 413 hardware.

414 6.2. Effect of Compute to I/O Ratio on Performance

415 The results in section 6.1 indicated opening the trajectory, communication, and read
 416 I/O to be important factors that appeared to correlate with stragglers. In order to better
 417 characterize the RMSD task, we computed the ratio between the time to complete the
 418 computation and the time spent on I/O per frame. The average values were $\overline{t_{\text{comp}}^{\text{frame}}} =$
 419 0.09 ms, $\overline{t_{\text{IO}}^{\text{frame}}} = 0.3$ ms, resulting in a compute-to-I/O ratio $R_{\text{comp/IO}} \approx 0.3$ (Eq. 7).
 420 Because $R_{\text{comp/IO}} \ll 1$, the RMSD analysis task was characterized as I/O bound.

421 As we were not able to achieve good scaling beyond a single node, we hypothesized
 422 that decreasing the I/O load relative to the compute load would interleave read I/O
 423 with longer periods of computation, thus reducing the impact of I/O contention and
 424 the impact of stragglers. We therefore set out to measure compute bound tasks, i.e.
 425 ones with $R_{\text{comp/IO}} \gg 1$. To measure the effect of the $R_{\text{comp/IO}}$ ratio on performance
 426 but leaving other parameters the same, we artificially increased the computational load
 427 by repeating the same RMSD calculation (line 10, algorithm 1) 40, 70 and 100 times
 428 in a loop, resulting in forty-fold (“40×”), seventy-fold (“70×”), and one hundred-fold
 429 (“100×”) load increases.

430 6.2.1. Effect of Increased Compute Workload

431 For an X -fold increase in workload, we expected the workload for the compu-
 432 tation to scale with X as $t_{\text{comp}}(X) = N_{\text{frames}}^{\text{total}} X \overline{t_{\text{comp}}^{\text{frame}}}$ while the read I/O workload
 433 $t_{\text{IO}}(X) = N_{\text{frames}}^{\text{total}} \overline{t_{\text{IO}}^{\text{frame}}}$ (number of frames times the average time to read a frame)
 434 should remain independent of X . Therefore, the ratio for any X should be $R_{\text{comp/IO}}(X) =$
 435 $t_{\text{comp}}(X)/t_{\text{IO}}(X) = X R_{\text{comp/IO}}(X = 1)$, i.e., $R_{\text{comp/IO}}$ should just linearly scale with the
 436 workload factor X . The measured $R_{\text{comp/IO}}$ ratios of 11, 19, 27 for the increased com-
 437 putational workloads agreed with this theoretical analysis, as shown in Table 4.

Workload X	t_{comp} (s)	$t_{\text{I/O}}$ (s)	$R_{\text{comp/I/O}}$	
			measured	theoretical
1×	226	791	0.29	
40×	8655	791	11	11
70×	15148	791	19	20
100×	21639	791	27	29

Table 4: Change in $R_{\text{comp/I/O}}$ ratio with change in the RMSD workload X . The RMSD workload was artificially increased in order to examine the effect of compute to I/O ratio on the performance. The reported compute and I/O time were measured based on the serial version using one core. The theoretical $R_{\text{comp/I/O}}$ (see text) is provided for comparison.




Figure 3: Effect of $R_{\text{comp/I/O}}$ ratio on performance of the RMSD task on *SDSC Comet*. We tested performance for $R_{\text{comp/I/O}}$ ratios of 0.3, 11, 19, 27, which correspond to 1× RMSD, 40× RMSD, 70× RMSD, and 100× RMSD respectively. (a) Effect of $R_{\text{comp/I/O}}$ on the speed-up. (b) Change in speed-up with respect to $R_{\text{comp/I/O}}$ for different processor counts. (c) Change in the efficiency with respect to $R_{\text{comp/I/O}}$ for different processor counts.

We performed the experiments with increased workload to measure the effect of the $R_{\text{comp/I/O}}$ ratio on performance (Figure 3). The strong scaling performance as measured by the speed-up $S(N)$ improved with increasing $R_{\text{comp/I/O}}$ ratio (Figure 3a). The calculations consistently showed better scaling performance to larger numbers of cores for higher $R_{\text{comp/I/O}}$ ratios, e.g., $N = 56$ cores for the 70× RMSD task. The speed-up and efficiency approached their ideal value for each processor count with increasing $R_{\text{comp/I/O}}$ ratio (Figures 3b and 3c). Even for moderately compute-bound workloads, such as the 40× and 70× RMSD tasks, increasing the computational workload over I/O reduced the impact of stragglers even though they still contributed to large variations in timing across different ranks and thus irregular scaling.

We also investigated the influence of the ratio of compute to communication costs

449 ($R_{\text{comp/comm}}$, Eq. 8) on performance in Appendix B. We found evidence to support the
 450 hypothesis that a larger ratio was beneficial, provided I/O costs could also be reduced.
 451 However, read I/O ultimately seemed to be the key determinant for performance, as
 452 discussed in the next sections.

453 6.2.2. Effect of Absence of Read I/O on Communication

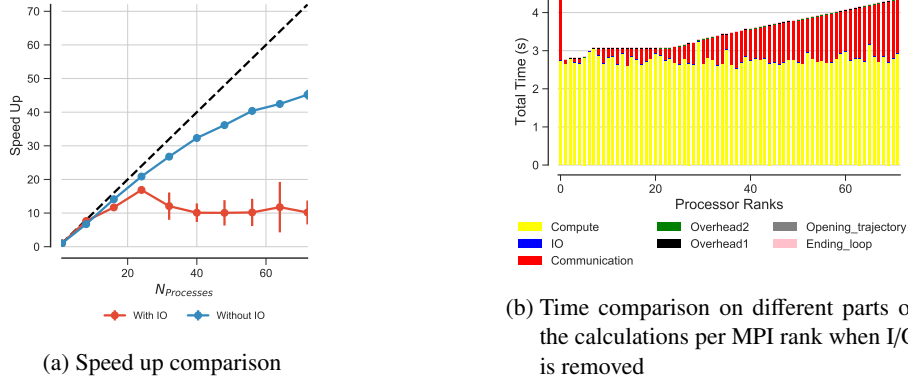


Figure 4: Comparison of the performance of the RMSD task with I/O ($R_{\text{comp/IO}} \approx 0.3$) and without I/O ($R_{\text{comp/IO}} = +\infty$) on *SDSC Comet*. Five repeats were performed to collect statistics. (a) Speed-up. The error bars show standard deviation with respect to the mean. (b) Compute t_{comp} , read I/O $t_{\text{I/O}} = 0$, communication t_{comm} , ending the for loop $t_{\text{end_loop}}$, opening the trajectory $t_{\text{opening_trajectory}}$, and overheads $t_{\text{overhead1}}$, $t_{\text{overhead2}}$ per MPI rank. (See Table 3 for definitions.)

454 In order to study an extreme case of a compute-bound task, we eliminated all I/O
 455 from the RMSD task by randomly generating artificial trajectory data in memory; the
 456 data had the the same size as if they had been obtained from the trajectory file. The
 457 time for the data generation was excluded and no file access was necessary. Without
 458 any I/O, performance improved markedly (Figure 4), with reasonable scaling up to
 459 72 cores (3 nodes). No stragglers were observed because overall communication time
 460 decreased and showed less variability; nevertheless, an increase in communication time
 461 prevented ideal scaling performance. Although in practice I/O cannot be avoided, this
 462 experiment demonstrated that accessing the trajectory file on the Lustre file system is
 463 at least one cause for the observed stragglers.

464 6.3. Reducing I/O Cost

465 In order to improve performance we needed to employ strategies to avoid the com-
466 petition over file access across different ranks when the $R_{\text{comp}/\text{IO}}$ ratio was small. To
467 this end, we experimented with two different ways for reducing the I/O cost: 1) split-
468 ting the trajectory file into as many segments as the number of processes, thus using
469 file-per-process access, and 2) using the HDF5 file format together with MPI-IO par-
470 allel reads instead of the XTC trajectory format. We discuss these two approaches and
471 their performance improvements in detail in the following sections.

472 6.3.1. Splitting the Trajectories (“subfiling”)

473 Subfiling is a mechanism previously used for splitting a large multi-dimensional
474 global array to a number of smaller subarrays in which each smaller array is saved in
475 a separate file. Subfiling reduces the file system control overhead by decreasing the
476 number of processes concurrently accessing a shared file [72, 73]. Because subfiling
477 is known to improve programming flexibility and performance of parallel shared-file
478 I/O, we investigated splitting our trajectory file into as many trajectory segments as
479 the number of processes. The trajectory file was split into N segments, one for each
480 process, with each segment having N_b frames. This way, each process would access its
481 own trajectory segment file without competing for file accesses.

482 We ran a benchmark up to 8 nodes (192 cores) and observed rather better scaling
483 behavior with efficiencies above 0.6 (Figure 5b and 5c) with the delay time for strag-
484 glers reduced from 65 s to about 10 s for 72 processes. However, scaling was still
485 far from ideal due to the MPI communication costs. Although the delay due to com-
486 munication was much smaller than compared to parallel RMSD with shared-file I/O
487 [compare Figure 5d ($t_{\text{comm}}^{\text{Straggler}} \gg t_{\text{comp}} + t_{\text{I/O}}$) to Figure 2d ($t_{\text{comm}}^{\text{Straggler}} \approx t_{\text{comp}} + t_{\text{I/O}}$)], it
488 was still delaying several processes and resulted in longer job completion times (Fig-
489 ure 5d). These delayed tasks impacted performance so that speed-up remained far from
490 ideal (Figure 5c).

491 The subfiling approach appeared promising but it required preprocessing of trajec-
492 tory files and additional storage space for the segments. We benchmarked the neces-
493 sary time for splitting the trajectory in parallel using different number of MPI processes

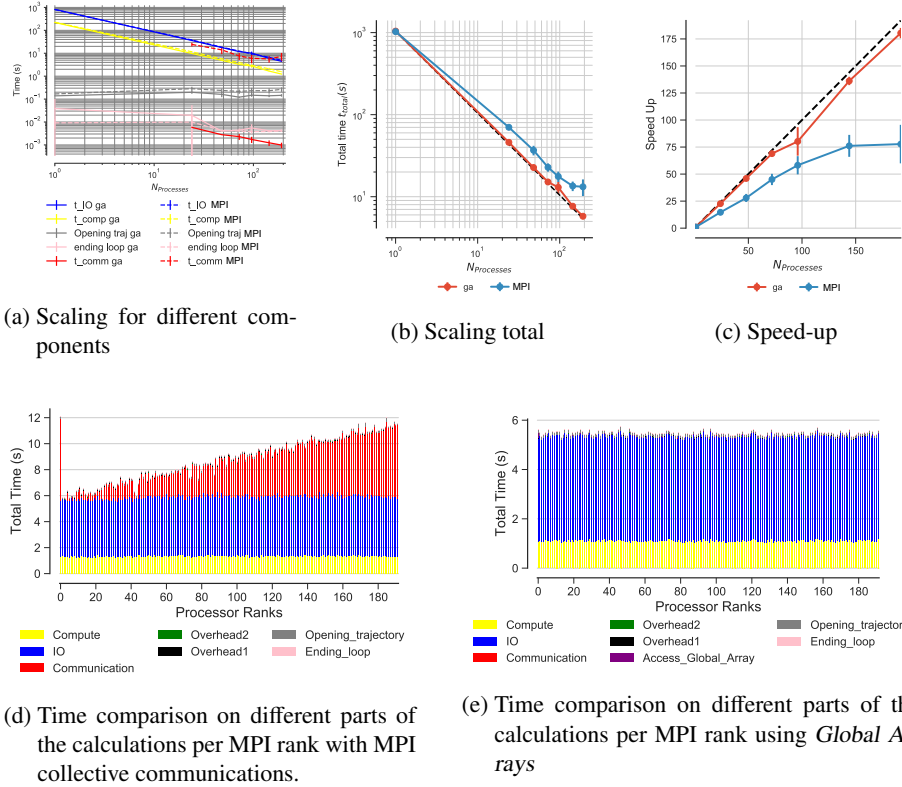


Figure 5: Comparison of the performance of the RMSD task on *SDSC Comet* when the trajectories are split (*subfiling*). The communication step used either MPI collective communications (“MPI”, with `MPI_Gather()`) or *Global Arrays* (“ga”, as described in Section 6.4). In the case of *Global Arrays*, all ranks updated the global RMSD array (`ga_put()`) and rank 0 accessed the whole RMSD array through the global memory address (`ga_get()`). Five repeats were performed to collect statistics. (a) Compute and I/O scaling versus number of processes. In serial, there is no communication and no data points are shown for $N = 1$. (b) Total time scaling versus number of processes. (c) Speed-up. (a-c) The error bars show standard deviation with respect to the mean. (d-e) Compute t_{comp} , read I/O t_{IO} , communication t_{comm} , access to the whole global RMSD array by rank 0, $t_{\text{Access_Global_Array}}$, ending the for loop $t_{\text{end_loop}}$, opening the trajectory $t_{\text{opening_trajectory}}$, and overheads $t_{\text{overhead1}}$, $t_{\text{overhead2}}$ per MPI rank; see Table 3 for the definitions.

494 (Table 5); in general the operation scaled well, with efficiencies > 0.8 although perfor-
 495 mance fluctuated, as seen for the case on six nodes where the efficiency dropped to
 496 0.34 for the run. These preprocessing times were not included in the estimates be-
 497 cause we focused on better understanding the principal causes of stragglers and we
 498 wanted to make the results directly comparable to the results of the previous sections.
 499 Nevertheless, from an end user perspective, preprocessing of trajectories can be inte-
 500 grated in workflows (especially as the data in Table 5 indicated good scaling) and the
 501 preprocessing time can be quickly amortized if the trajectories are analyzed repeatedly.

N_{seg}	N_p	nodes	time (s)	S	E
24	24	1	89.9	1.0	1.0
48	48	2	46.8	1.9	0.96
72	72	3	33.7	2.7	0.89
96	96	4	25.1	3.6	0.89
144	144	6	43.7	2.1	0.34
192	192	8	13.5	6.7	0.83


Table 5: The wall-clock time spent for writing N_{seg} trajectory segments using N_p processes using MPI on *SDSC Comet*. One set of runs was performed for the timings. Scaling S and efficiency E are relative to the 1 node case (24 MPI processes).

Often trajectories from MD simulations on HPC machines are produced and kept in smaller files or segments that can be concatenated to form a full continuous trajectory file. These trajectory segments could be used for the subfiling approach. However, it might not be feasible to have exactly one segment per MPI rank, with all segments of equal size, which constitutes the ideal case for subfiling. MDAnalysis can create virtual trajectories from separate trajectory segment files that appear to the user as a single trajectory. In Appendix C we investigated if this so-called *ChainReader* functionality could benefit from the subfiling approach. We found some improvements in performance but discovered limitations in the design of the ChainReader (namely that all segment files are initially opened) that will have to be addressed before equivalent performance can be reached.


6.3.2. MPI-based Parallel HDF5

In the HPC community, parallel I/O with MPI-IO is widely used in order to address I/O limitations. We investigated MPI-based Parallel HDF5 to improve I/O scaling. We converted our XTC trajectory file into a simple custom HDF5 format so that we could test the performance of parallel I/O with the HDF5 file format. The code for this file format conversion can be found in the GitHub repository. The time it took to convert our XTC file with 2,512,200 frames into HDF5 format was about 5,400 s on a local workstation with network file system (NFS).


We ran our benchmark on up to 16 nodes (384 cores) on *SDSC Comet* and we observed near ideal scaling behavior (Figures 6a and 6b) with parallel efficiencies above 0.8 on up to 8 nodes (Figure A.13a) with no straggler tasks (Figure 6d). The trajec-




(a) Scaling total



(b) Speed-up



(c) Scaling for different components



(d) Time comparison on different parts of the calculations per MPI rank

Figure 6: Performance of the RMSD task with MPI-based parallel HDF5 (MPI-IO) on *SDSC Comet*. Data are read from the file system from a shared HDF5 file format instead of XTC format (independent I/O) and results are communicated back to rank 0. Five repeats were performed to collect statistics. (a-c) The error bars show standard deviation with respect to the mean. In serial, there is no communication and no data points are shown for $N = 1$ in (c). (d) Compute t_{comp} , read I/O $t_{\text{I/O}}$, communication t_{comm} , ending the for loop $t_{\text{end_loop}}$, opening the trajectory $t_{\text{opening_trajectory}}$, and overheads $t_{\text{overhead1}}$, $t_{\text{overhead2}}$ per MPI rank; see Table 3 for definitions. These are typical data from one run of the five repeats.

524 tory reading I/O ($t_{\text{I/O}}$) was the dominant contribution, followed by compute (t_{comp}), but
 525 because both contributions scaled well, overall scaling performance remained good,
 526 even for 384 cores. We observed a low-performing outlier for 12 nodes (288 cores)
 527 with slower I/O than typical but did not further investigate. Importantly, the trajec-
 528 tory opening cost remained negligible (in the millisecond range) and the cost for MPI

communication also remained small (below 0.1 s, even for 16 nodes). Overall, parallel MPI with HDF5 appeared to be a robust approach to obtain good scaling, even for I/O-bound tasks.


6.4. Testing the Global Arrays Toolkit

The *Global Arrays* (GA) toolkit [33] is a convenient layer to represent and access arrays across multiple MPI ranks and nodes. Because of its convenience and possibly reduced communications overhead due to its use of shared memory on a physical node and MPI for inter-node communication (see Section 3.3) we wanted to compare parallel trajectory analysis with GA to the MPI-based implementation that was discussed in the previous sections.


With GA, one large RMSD array called the *global array* was defined and each MPI rank updated its associated block in the global RMSD array using `ga_put()` (Algorithm 2). At the end, when all the processes exited the `block_rmsd()` function and updated their local block in the global array, rank 0 accessed the whole global array using `ga_access()`. In GA, the time for communication is $t_{\text{ga_put}()} + t_{\text{ga_access}()}$. We tested that the approach with GA (Algorithm 2) gave the same results as the previously discussed approach with `MPI_Gather()` (Algorithm 1).

Shared file. Using GA improved the strong scaling performance (Figures 7a and 7b) by reducing the communication time. Nevertheless, the remaining variation in the trajectory I/O part of the calculation and in particular the initial opening of the trajectory prevented ideal scaling (Figure 7c). Stragglers were primarily due to the fact that all ranks had to open the same trajectory file at the beginning of the execution (Figure 7d). In this case, these slow processes took about 50 s, which was slower than the mean execution time of all other ranks of 17 s. Trajectory opening was already problematic in the initial test (Figure 2c), which was still dominated by the communication cost. By substantially reducing communication cost, the simultaneous trajectory opening by multiple ranks emerged as the next dominant cause for stragglers.


Subfiling. We tested subfiling (see Section 6.3.1) with GA to reduce the initial delay due to trajectory opening. Under otherwise identical conditions as in the previous sec-




(a) Scaling total



(b) Speed-up



(c) Scaling for different components



(d) Time comparison on different parts of the calculations per MPI rank

Figure 7: Performance of the RMSD task using *Global Arrays* on *SDSC Comet*. All ranks updated the global RMSD array (`ga_put()`) and rank 0 accessed the whole RMSD array through the global memory address (`ga_get()`). Five repeats were performed to collect statistics. (a-c) The error bars show standard deviation with respect to the mean. In serial, there is no communication and not data points are shown for $N = 1$ in (c). In (d), compute t_{comp} , read I/O $t_{\text{I/O}}$, communication t_{comm} , access to the whole global array by rank 0, $t_{\text{Access_Global_Array}}$, ending the for loop $t_{\text{end_loop}}$, opening the trajectory $t_{\text{opening_trajectory}}$, and overheads $t_{\text{overhead1}}$, $t_{\text{overhead2}}$ per MPI rank are shown; see Table 3 for definitions. These are typical data from one run of the five repeats. MPI ranks 20 and 56 were stragglers.

tion we now observed near ideal scaling behavior with efficiencies above 0.9 (Figure 5b
and 5c) without any straggler tasks (Figure 5e). Although the reason why in our case
GA appeared to be more efficient than direct MPI-based communication remained un-
clear, these results showed that contention for file access clearly impacted performance.

562 By removing the contention, near ideal scaling could be achieved.

563 6.5. Likely Causes of Stragglers

564 The data indicated that an increase in the duration of both MPI communication
565 and trajectory file access lead to large variability in the run time of individual MPI
566 processes and ultimately poor scaling performance beyond a single node. A discussion
567 of likely causes for stragglers begins with the observation that opening and reading a
568 single trajectory file from multiple MPI processes appeared to be at the center of the
569 problem.

570 In MDAnalysis, individual trajectory frames are loaded into memory, which en-
571 sures that even systems with tens of millions of atoms can be efficiently analyzed on re-
572 sources with moderate RAM sizes. The test trajectory (file size 30 GB) had 2,512,200
573 frames in total so each frame was about 0.011 MB in size. With $t_{I/O} \approx 0.3$ ms per
574 frame, the data were ingested at a rate of about 40 MB/s for a single process. For
575 24 MPI ranks (one SDSC Comet node), the aggregated reading rate would have been
576 about 1 GB/s, well within the available bandwidth of 56 Gb/s of the InfiniBand net-
577 work interface that served the Lustre file system, but nevertheless sufficient to produce
578 substantial constant network traffic.

579 Furthermore, in our study the default Lustre stripe size value was 1 MB, i.e., the
580 amount of contiguous data stored on a single Lustre object storage target (OST). Each
581 I/O request read a single Lustre stripe but because the I/O size (0.011 MB) was smaller
582 than the stripe size, many of these I/O requests were likely just accessing the same
583 stripe on the same OST but nevertheless had to acquire a new reading lock for each
584 request. The reason for this behavior is related to ensuring POSIX consistency that re-
585 lates to POSIX I/O API and POSIX I/O semantics, which can have adverse effects on
586 scalability and performance. Parallel file systems like Lustre implement sophisticated
587 distributed locking mechanisms to ensure consistency. For example, locking mecha-
588 nisms ensures that a node can not read from a file or part of a file while it might be being
589 modified by another node. In fact, when the application I/O is not designed in a way to
590 avoid scenarios where multiple nodes are fighting over locks for overlapping extents,
591 Lustre can suffer from scalability limitations [74]. Continuously keeping metadata up-

dated in order to have fully consistent reads and writes (POSIX metadata management), requires writing a new value for the file’s last-accessed time (POSIX atime) every time a file is read, imposing a significant burden on parallel file system [75]. It was observed that contention for the interconnect between OSTs and compute nodes due to MPI communication may lead to variable performance in I/O measurements [76]. Conversely, our data suggest that single-shared-file I/O on Lustre can negatively affect MPI communication as well, even at moderate numbers (tens to hundreds) of concurrent requests, similar to recent network simulations that predicted interference between MPI and I/O traffic [77]. This work indicated that MPI traffic (inter-process communication) can be affected by increasing I/O, and in particular, a few MPI processes were always delayed by one to two orders of magnitude more than the median time. In summary, these observations in the context of the work by Brown et al. [77] suggest that our observed stragglers with large variance in the communication step might be due to interference of MPI communications with the I/O requests on the same network.

7. Reproducibility and Performance Comparison on Different Clusters

In this section we compare the performance of the RMSD task on different HPC resources (Table 1) to examine the robustness of the methods we used for our performance study and to ensure that the results are general and independent from the specific HPC system. Scripts and instructions to set up the computational environments and reproduce our computational experiments are provided in a git repository as described in section 5.

In Appendix A, we demonstrated that stragglers occur on *PSC Bridges* (Figure A.11) and *LSU SuperMIC* (Figure A.12) in a manner similar to the one observed on *SDSC Comet* (section 6.1). We performed additional comparisons for several cases discussed previously, namely (1) splitting the trajectories with blocking collective communications in MPI, (2) splitting the trajectories with Global Arrays for communications, and (3) MPI-based parallel HDF5.

619 7.1. Splitting the Trajectories

620 Figure 8 shows the strong scaling of the RMSD task on different HPC resources.
 621 Splitting the trajectories with Global Arrays for communication resulted in very good
 622 scaling performance on *LSU SuperMIC*, similar to the results obtained on *SDSC Comet*.
 623 The results with MPI blocking collective communication (instead of Global Arrays)
 624 were also comparable between the two clusters, with scaling far from ideal due to the
 625 communication cost (see section 6.3.1 and Figures 5d and A.14). Overall, the scaling of
 626 the RMSD task is better on *LSU SuperMIC* than on *SDSC Comet* and the performance
 627 gap increased with increasing core number. The results on *LSU SuperMIC* confirmed
 628 the conclusion obtained on *SDSC Comet* that at least in this case Global Arrays per-
 629 formed better than MPI blocking collective communication.

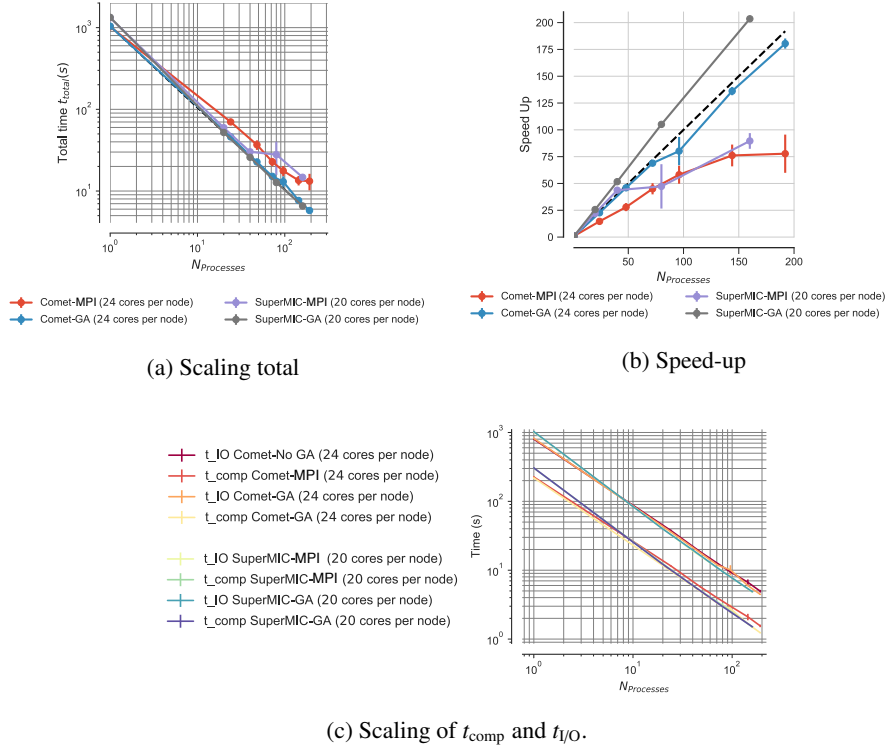


Figure 8: Comparison of the performance of the RMSD task across different clusters (*SDSC Comet*, *LSU SuperMIC*) when the trajectories are split (*subfiling*). Results were communicated back to rank 0 either with MPI collective communications (label “MPI”) or using *Global Arrays* (label “GA”). Five repeats were performed to collect statistics. The error bars show the standard deviation with respect to the mean.

630 7.2. MPI-based Parallel HDF5

631 Figure 9 shows the scaling on *SDSC Comet*, *LSU SuperMIC*, and *PSC Bridges*
632 using MPI-based parallel HDF5. Performance on *SDSC Comet* and *LSU SuperMIC*
633 was very good with near ideal linear strong scaling. The performance on *PSC Bridges*
634 was sensitive to how many cores per node were used. Using all 28 cores in a node
635 resulted in poor performance but decreasing the number of cores per node and equally
636 distributing processes over nodes improved the scaling (Figure 9), mainly by reducing
637 variation in the I/O times.

638 The main difference between the runs on *PSC Bridges* and *SDSC Comet/LSU Su-*
639 *perMIC* appeared to be the variance in $t_{I/O}$ (Figure 9c). The I/O time distribution was
640 fairly small and uniform across all ranks on *SDSC Comet* and *LSU SuperMIC* (Figures
641 10b and 6d). However, on *PSC Bridges* the I/O time was on average about two and a
642 half times larger and the I/O time distribution was also more variable across different
643 ranks (Figure 10a).

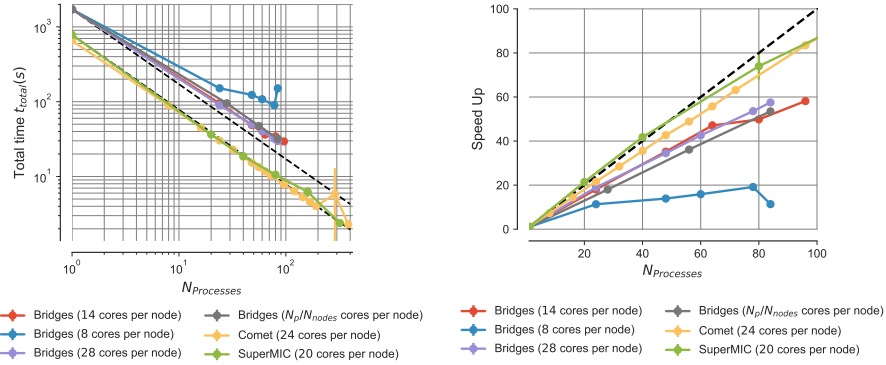
644 7.3. Comparison of Compute and I/O Scaling Across Different Clusters

645 A full comparison of compute and I/O scaling across different clusters for differ-
646 ent test cases and algorithms is shown in Table 6. For MPI-based parallel HDF5, both
647 the compute and I/O time on *Bridges* were consistently larger than their corresponding
648 values on *SDSC Comet* and *LSU SuperMIC*. For example, with one core the corre-
649 sponding compute and I/O time were $t_{\text{comp}} = 387$ s, $t_{I/O} = 1318$ s versus 225 s, 423 s
650 on *SDSC Comet* and 273 s, 503 s on *LSU SuperMIC*. This performance difference be-
651 came larger with increasing core number. When the trajectories were split and Global
652 Arrays was used for communication both *SDSC Comet* and *LSU SuperMIC* showed
653 similar performance.

654 Overall, the results from *SDSC Comet* and *LSU SuperMIC* are consistent with each
655 other. Performance on *PSC Bridges* seemed sensitive to the exact allocation of cores on
656 each node but nevertheless the approaches that decreased the occurrence of stragglers
657 on *SDSC Comet* and *LSU SuperMIC* also improved performance on *PSC Bridges*. Thus,
658 the findings described in the previous sections are valid for a range of different HPC
659 clusters with Lustre file systems.


Cluster	Gather	File Access	Time	Serial	$N_{Processes}$									
					Comet: 24	Comet: 48	Comet: 72	Comet: 96	Comet: 144	Comet: 192	Comet: 384	SuperMIC: 20	SuperMIC: 40	SuperMIC: 80
Comet	MPI	Single	t_{IO}	791 \pm 5.22	49 \pm 3.45	29 \pm 1.3	26 \pm 9.19	–	–	–	–	–	–	–
			t_{comp}	225 \pm 5.4	11 \pm 0.75	6 \pm 0.35	4 \pm 0.48	–	–	–	–	–	–	–
Bridges	MPI	Single	t_{IO}	770 \pm 10.8	38 \pm 0.84	33 \pm 19.4	15 \pm 1.6	–	–	–	–	–	–	–
			t_{comp}	221 \pm 3.9	11 \pm 0.43	6 \pm 0.32	4 \pm 0.18	–	–	–	–	–	–	–
SuperMIC	MPI	Single	t_{IO}	1014.51 \pm 2.94	48.08 \pm 0.35	24.5 \pm 0.79	12 \pm 0.31	–	6.24 \pm 0.38	–	–	–	–	–
			t_{comp}	303.85 \pm 2.3	14.56 \pm 0.14	7.4 \pm 0.25	3.7 \pm 0.12	–	1.8 \pm 0.04	–	–	–	–	–
Comet	GA	Single	t_{IO}	820 \pm 18.49	41 \pm 8.99	23 \pm 4.14	15 \pm 2.06	–	–	–	–	–	–	–
			t_{comp}	219 \pm 9.8	10 \pm 0.3	5 \pm 0.48	3 \pm 0.54	–	–	–	–	–	–	–
Comet	MPI	Splitting	t_{IO}	799 \pm 5.22	37 \pm 1.22	18 \pm 0.18	12 \pm 0.14	9 \pm 0.3	6 \pm 0.66	4 \pm 0.23	–	–	–	–
			t_{comp}	225 \pm 5.4	11 \pm 0.31	5 \pm 0.07	3 \pm 0.04	3 \pm 0.11	2 \pm 0.23	1 \pm 0.07	–	–	–	–
SuperMIC	MPI	Splitting	t_{IO}	1013.75 \pm 2.8	39.99 \pm 0.36	19.18 \pm 0.25	9.61 \pm 0.28	–	4.83 \pm 0.06	–	–	–	–	–
			t_{comp}	304.26 \pm 2.55	12.41 \pm 0.22	5.99 \pm 0.09	3.08 \pm 0.13	–	1.5 \pm 0.01	–	–	–	–	–
Comet	GA	Splitting	t_{IO}	820 \pm 18.5	36 \pm 0.78	17 \pm 0.3	11 \pm 0.23	10 \pm 1.7	5 \pm 0.14	4 \pm 0.07	–	–	–	–
			t_{comp}	219 \pm 9.5	9 \pm 0.22	4 \pm 0.07	3 \pm 0.04	2 \pm 0.4	1 \pm 0.05	1 \pm 0.02	–	–	–	–
SuperMIC	GA	Splitting	t_{IO}	1027.62 \pm 10.32	39.62 \pm 0.2	19.66 \pm 0.1	9.57 \pm 0.1	–	4.86 \pm 0.05	–	–	–	–	–
			t_{comp}	305.78 \pm 3.47	12.16 \pm 0.1	6.01 \pm 0.007	2.97 \pm 0.1	–	1.51 \pm 0.03	–	–	–	–	–
Comet	MPI	PHDF5	t_{IO}	423 \pm 5.88	19 \pm 0.3	9 \pm 0.13	6 \pm 0.06	5 \pm 0.12	3 \pm 0.2	3 \pm 0.25	1.57 \pm 0.29	–	–	–
			t_{comp}	225 \pm 6.55	10 \pm 0.12	5 \pm 0.1	3 \pm 0.04	2 \pm 0.05	1 \pm 0.04	1 \pm 0.03	0.76 \pm 0.09	–	–	–
Bridges	MPI	PHDF5	t_{IO}	1318.87 \pm 10.42	67.93 \pm 0.52	37.37 \pm 0.2	30.35 \pm 0.15	24.16 \pm 0.89	22.5 \pm 0.17	–	–	–	–	–
			t_{comp}	387.8 \pm 5.51	21.97 \pm 0.38	12.12 \pm 0.34	9.79 \pm 0.24	7.72 \pm 0.03	7.18 \pm 0.08	–	–	–	–	–
SuperMIC	MPI	PHDF5	t_{IO}	503.69 \pm 2.57	12.96 \pm 0.06	6.46 \pm 0.02	3.2 \pm 0.01	–	1.64 \pm 0.01	–	0.82 \pm 0.004	–	–	–
			t_{comp}	273.54 \pm 4.7	23.44 \pm 0.29	12.22 \pm 0.43	7.3 \pm 0.85	–	4.59 \pm 0.96	–	1.55 \pm 0.009	–	–	–

Table 6: Comparison of the compute and I/O scaling for different test cases and number of processes. Five repeats were performed to collect statistics. The mean value and the standard deviation with respect to mean are reported for each case.



(a) Scaling total

(b) Speed-up



(c) Scaling of t_{comp} and t_{IO}

Figure 9: Comparison of the performance of the RMSD task across different clusters (*SDSC Comet*, *PSC Bridges*, *LSU SuperMIC*) with MPI-IO. Data were read from a shared HDF5 file instead of an XTC file, using MPI independent I/O in the PHDF5 library. Results were communicated back to rank 0. N_p/N_{nodes} indicates that number of processes used for the task were equally distributed over all compute nodes. Five repeats were performed to collect statistics. The error bars show standard deviation with respect to mean. In (b) only results up to 100 cores are shown to simplify the comparison; see Fig. 6b for *SDSC Comet* and Fig. A.13c for *LSU SuperMic* data.

8. Guidelines for Improving Parallel Trajectory Analysis Performance

Although the performance measurements were performed with *MDAnalysis* and therefore capture some details of this library such as the specific timings for file reading, we believe that the broad picture is fairly general and applies to any Python-based approach that uses MPI for parallelizing trajectory access with a split-apply-combine approach. Based on the lessons that we learned, we suggest the following guidelines

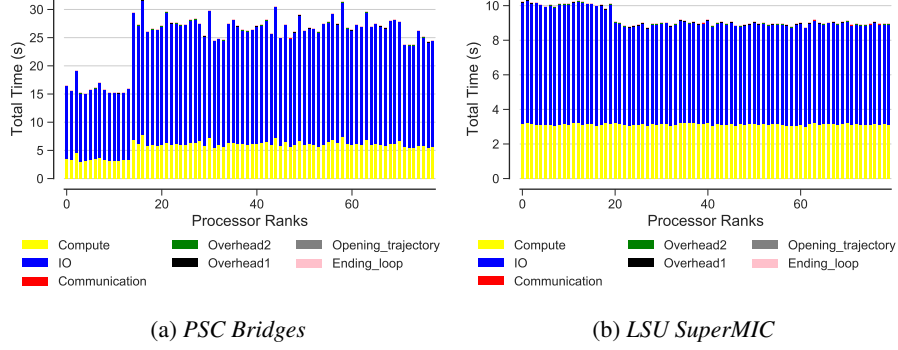


Figure 10: Examples of timing per MPI rank for the RMSD task with MPI-based parallel HDF5 on (a) *PSC Bridges* and (b) *LSU SuperMIC*. Five repeats were performed to collect statistics and these were typical data from one run of the five repeats. Compute t_{comp} , read I/O $t_{\text{I/O}}$, communication t_{comm} , ending the for loop $t_{\text{end_loop}}$, opening the trajectory $t_{\text{opening_trajectory}}$, and overheads $t_{\text{overhead1}}$, $t_{\text{overhead2}}$ per MPI rank; see Table 3 for definitions.

to improve strong scaling performance:

Heuristic 1 Calculate compute to I/O ratio ($R_{\text{comp}/\text{IO}}$, Eq. 7) and compute to communication ratio ($R_{\text{comp}/\text{comm}}$, Eq. 8). $R_{\text{comp}/\text{IO}}$ determines whether the task is compute bound ($R_{\text{comp}/\text{IO}} \gg 1$) or IO bound ($R_{\text{comp}/\text{IO}} \ll 1$). $R_{\text{comp}/\text{comm}}$ determines whether the task is communication bound ($\frac{t_{\text{comp}}}{t_{\text{comm}}} \ll 1$) or compute bound ($R_{\text{comp}/\text{IO}} \gg 1$).

As discussed in Section 6.2, for I/O bound problems the interference between MPI communication and I/O traffic can be problematic [50, 77] and the performance of the task will be affected by the straggler tasks that delay job completion time.

Heuristic 2 For $R_{\text{comp}/\text{IO}} \geq 1$, single-shared-file I/O can be used and will not decrease performance. One may consider the following cases:

Heuristic 2.1 If $R_{\text{comp}/\text{comm}} \gg 1$, the task is compute bound and will scale well as shown in Figure 3.

Heuristic 2.2 If $R_{\text{comp}/\text{comm}} \ll 1$, one might consider using *Global Arrays* to improve scaling by utilizing efficient distribution of data via the shared arrays (section 6.4).

Heuristic 3 For $R_{\text{comp}/\text{IO}} \leq 1$ the task is I/O bound and single-shared-file I/O should be avoided to remove unnecessary metadata operations. One might want to consider

683 the following steps:

684 **Heuristic 3.1** If there is access to HDF5 format, use MPI-based Parallel HDF5 (Sec-
685 tion 6.3.2).

686 **Heuristic 3.2** If the trajectory file is not in HDF5 format then one can consider
687 subfilig and split the single trajectory file into as many trajectory segments
688 as the number of processes. Splitting the trajectories can be easily performed
689 in parallel and trajectory conversion may be integrated into the beginning of
690 standard workflows for MD simulations. Alternatively, trajectories may be
691 kept in smaller chunks if they are already produced in batches; for instance,
692 when running simulations with *Gromacs* [55], the `gmx mdrun -noappend`
693 option produces individual trajectory segments instead of extending an existing
694 trajectory file.

695 **Heuristic 3.3** In case of $R_{\text{comp/comm}} \ll 1$, use of *Global Arrays* may be considered
696 to potentially improve scaling (Section 6.3.1).

697 9. Conclusions

698 We analyzed the strong scaling performance of a typical task when analysing MD
699 trajectories, the calculation of the time series of the RMSD of a protein, with the widely
700 used Python-based *MDAnalysis* library. The task was parallelized with MPI following
701 the *split-apply-combine* approach by having each MPI process analyze a contiguous
702 segment of the trajectory. This approach did not scale beyond a single node because
703 straggler MPI processes exhibited large upward variations in runtime. Stragglers were
704 primarily caused by either increased MPI communication costs or increased time to
705 open the single shared trajectory file whereas both the computation and the ingestion
706 of data exhibited close to ideal strong scaling behavior. Stragglers were less prevalent
707 for compute-bound workloads (i.e., $R_{\text{comp/IO}} \gg 1$), suggesting that file read I/O was
708 responsible for poor MPI communication. In particular, artificially removing all I/O
709 substantially improved performance of the communication step and thus brought over-
710 all performance close to ideal (i.e., linear increase in speed-up with processor count

711 with slope one). By performing benchmarks on three different XSEDE supercomput-
712 ers we showed that our results were independent from the specifics of the hardware
713 and local environment. Our results hint at the possibility that stragglers might be due
714 to the competition between MPI messages and the Lustre file system on the shared In-
715 finiBand interconnect, which would be consistent with other similar observations [50]
716 and theoretical predictions by Brown et al. [77]. One possible interpretation of our re-
717 sults is that for a sufficiently large per-frame compute workload, read I/O interferes
718 much less with communication than for an I/O bound task that almost continuously ac-
719 cesses the file system. This interpretation suggested that we needed to improve read I/O
720 to reduce interference.

721 We investigated subfilng (splitting of the trajectories into separate files, one for
722 each MPI rank) and MPI-based parallel I/O. Subfilng improved scaling, especially
723 when combined with the *Global Arrays* toolkit. *Global Arrays* reduced the commu-
724 nication cost compared to MPI collective communications even though it only acts as
725 programming layer to access data across multiple nodes in a convenient array form and
726 also uses MPI for its inter-node data exchange. Subfilng with *Global Arrays* achieved
727 nearly ideal scaling up to 192 cores (8 nodes on *SDSC Comet*). When we used MPI-
728 based parallel I/O through HDF5 together with MPI for communications we achieved
729 nearly ideal performance up to 384 cores (16 nodes on *SDSC Comet*) and speed-ups of
730 two orders of magnitude compared to the serial execution. The latter approach appears
731 to be a promising way forward as it directly builds on very widely used technology
732 (MPI-IO and HDF5) and echoes the experience of the wider HPC community that par-
733 allel file I/O is necessary for efficient data handling.

734 The biomolecular simulation community suffers from a large number of trajectory
735 file formats with very few being based on HDF5, with the exception of the H5MD
736 format [78] and the MDTraj HDF5 format [20]. Our work suggests that HDF5-based
737 formats should be seriously considered as the default for MD simulations if users want
738 to make efficient use of their HPC systems for analysis. Alternatively, enabling MPI-
739 IO for trajectory readers in libraries such as *MDAnalysis* might also provide a path
740 forward to better read performance.

741 We summarized our findings in a number of guidelines for improving the scaling

of parallel analysis of MD trajectory data. We showed that it is feasible to run an I/O bound analysis task on HPC resources with a Lustre parallel file system and achieve good scaling behavior up to 384 CPU cores with an almost 300-fold speed-up compared to serial execution. Although we focused on the *MDAnalysis* library, similar strategies are likely to be more generally applicable and useful to the wider biomolecular simulation community.

Acknowledgements

We are grateful to Sarp Oral for insightful comments on this manuscript. This work was supported by the National Science Foundation under grant numbers ACI-1443054 and ACI-1440677. This work used the Extreme Science and Engineering Discovery Environment (XSEDE), which is supported by National Science Foundation grant number ACI-1548562. *SDSC Comet* at the San Diego Supercomputer Center, *LSU SuperMic* at Louisiana State University, and *PSC Bridges* at the Pittsburgh Supercomputing Center were used under allocations TG-MCB090174 and TG-MCB130177.

References

1. D. W. Borhani, D. E. Shaw, The future of molecular dynamics simulations in drug discovery, *J Comput Aided Mol Des* 26 (2012) 15–26. doi:10.1007/s10822-011-9517-y.
2. R. O. Dror, R. M. Dirks, J. P. Grossman, H. Xu, D. E. Shaw, Biomolecular simulation: a computational microscope for molecular biology, *Annu Rev Biophys* 41 (2012) 429–52. doi:10.1146/annurev-biophys-042910-155245.
3. M. Orozco, A theoretical view of protein dynamics, *Chem. Soc. Rev.* 43 (2014) 5051–5066. doi:10.1039/C3CS60474H.
4. J. R. Perilla, B. C. Goh, C. K. Cassidy, B. Liu, R. C. Bernardi, T. Rudack, H. Yu, Z. Wu, K. Schulten, Molecular dynamics simulations of large macromolecular complexes, *Current Opinion in Structural Biology* 31 (2015) 64–74. URL: <http://www.sciencedirect.com/science/article/pii/S0959440X15000342>. doi:<http://dx.doi.org/10.1016/j.sbi.2015.03.007>.

- 771 5. S. Bottaro, K. Lindorff-Larsen, Biophysical experiments and biomolecular sim-
772 ulations: A perfect match?, *Science* 361 (2018) 355–360. doi:10.1126/
773 science.aat4010.
- 774 6. M. E. Tuckerman, *Statistical Mechanics: Theory and Molecular Simulation*,
775 Oxford University Press, Oxford, UK, 2010.
- 776 7. C. Mura, C. E. McAnany, An introduction to biomolecular simulations and
777 docking, *Molecular Simulation* 40 (2014) 732–764. URL: <http://dx.doi.org/10.1080/08927022.2014.935372>. doi:10.1080/08927022.
778 2014.935372.
- 779 8. T. Cheatham, D. Roe, The impact of heterogeneous computing on workflows
780 for biomolecular simulation and analysis, *Computing in Science Engineering*
781 17 (2015) 30–39. doi:10.1109/MCSE.2015.7.
- 782 9. G. R. Kneller, V. Keiner, M. Kneller, M. Schiller, nmoldyn: A
783 program package for a neutron scattering oriented analysis of molecu-
784 lar dynamics simulations, *Computer Physics Communications* 91 (1995)
785 191 – 214. URL: [http://www.sciencedirect.com/science/](http://www.sciencedirect.com/science/article/pii/001046559500048K)
786 [article/pii/001046559500048K](http://www.sciencedirect.com/science/article/pii/001046559500048K). doi:[http://dx.doi.org/10.](http://dx.doi.org/10.1016/0010-4655(95)00048-K)
787 [1016/0010-4655\(95\)00048-K](http://dx.doi.org/10.1016/0010-4655(95)00048-K).
- 788 10. K. Hinsén, E. Pellegrini, S. Stachura, G. R. Kneller, nmoldyn 3: Us-
789 ing task farming for a parallel spectroscopy-oriented analysis of molec-
790 ular dynamics simulations, *Journal of Computational Chemistry* 33
791 (2012) 2043–2048. URL: <http://dx.doi.org/10.1002/jcc.23035>.
792 doi:10.1002/jcc.23035.
- 793 11. W. Humphrey, A. Dalke, K. Schulten, VMD – Visual Molecular Dynam-
794 ics, *J. Mol. Graph.* 14 (1996) 33–38. URL: [http://www.ks.uiuc.edu/](http://www.ks.uiuc.edu/Research/vmd/)
795 [Research/vmd/](http://www.ks.uiuc.edu/Research/vmd/).
- 796 12. K. Hinsén, The molecular modeling toolkit: a new approach to molecular sim-
797 ulations, *Journal of Computational Chemistry* 21 (2000) 79–85.
- 798 13. B. J. Grant, A. P. C. Rodrigues, K. M. ElSawy, J. A. McCammon, L. S. D. Caves,
799 Bio3d: an r package for the comparative analysis of protein structures, *Bioin-*
800 *formatics* 22 (2006) 2695–6. doi:10.1093/bioinformatics/btl1461.
- 801

- 802 14. T. Tu, C. A. Rendleman, D. W. Borhani, R. O. Dror, J. Gullingsrud, M. O.
803 Jensen, J. L. Klepeis, P. Maragakis, P. Miller, K. A. Stafford, D. E. Shaw, A
804 scalable parallel framework for analyzing terascale molecular dynamics sim-
805 ulation trajectories, in: 2008 SC - International Conference for High Per-
806 formance Computing, Networking, Storage and Analysis, 2008, pp. 1–12.
807 doi:10.1109/SC.2008.5214715.
- 808 15. T. D. Romo, A. Grossfield, LOOS: An extensible platform for the structural
809 analysis of simulations, in: 31st Annual International Conference of the IEEE
810 EMBS, IEEE, Minneapolis, Minnesota, USA, 2009, pp. 2332–2335.
- 811 16. T. D. Romo, N. Leioatts, A. Grossfield, Lightweight object oriented structure
812 analysis: Tools for building tools to analyze molecular dynamics simulations,
813 Journal of Computational Chemistry 35 (2014) 2305–2318. URL: [http://](http://dx.doi.org/10.1002/jcc.23753)
814 dx.doi.org/10.1002/jcc.23753. doi:10.1002/jcc.23753.
- 815 17. N. Michaud-Agrawal, E. J. Denning, T. B. Woolf, O. Beckstein, MDAnalysis:
816 A toolkit for the analysis of molecular dynamics simulations, J Comp Chem 32
817 (2011) 2319–2327. doi:10.1002/jcc.21787.
- 818 18. R. J. Gowers, M. Linke, J. Barnoud, T. J. E. Reddy, M. N. Melo, S. L.
819 Seyler, D. L. Dotson, J. Domański, S. Buchoux, I. M. Kenney, O. Beckstein,
820 MDAnalysis: A Python package for the rapid analysis of molecular dynam-
821 ics simulations, in: S. Benthall, S. Rostrup (Eds.), Proceedings of the 15th
822 Python in Science Conference, SciPy, Austin, TX, 2016, pp. 102 – 109. URL:
823 <http://mdanalysis.org>.
- 824 19. D. R. Roe, I. Thomas E. Cheatham, Ptraj and cpptraj: Software for processing
825 and analysis of molecular dynamics trajectory data, Journal of Chemical Theory
826 and Computation 9 (2013) 3084–3095. URL: [http://dx.doi.org/10.](http://dx.doi.org/10.1021/ct400341p)
827 [1021/ct400341p](http://dx.doi.org/10.1021/ct400341p). doi:10.1021/ct400341p, PMID: 26583988.
- 828 20. R. T. McGibbon, K. A. Beauchamp, M. P. Harrigan, C. Klein, J. M.
829 Swails, C. X. Hernández, C. R. Schwantes, L.-P. Wang, T. J. Lane, V. S.
830 Pande, MDTraj: A modern open library for the analysis of molecular dy-
831 namics trajectories, Biophysical Journal 109 (2015) 1528 – 1532. URL:
832 <http://www.sciencedirect.com/science/article/pii/>

- 833 S0006349515008267. doi:10.1016/j.bpj.2015.08.015.
- 834 21. S. O. Yesylevskyy, Pteros 2.0: Evolution of the fast parallel molecular anal-
 835 ysis library for c++ and python, *Journal of Computational Chemistry* 36
 836 (2015) 1480–1488. URL: <http://dx.doi.org/10.1002/jcc.23943>.
 837 doi:10.1002/jcc.23943.
- 838 22. S. Doerr, M. J. Harvey, F. Noé, G. De Fabritiis, HTMD: High-throughput molec-
 839 ular dynamics for molecular discovery, *Journal of Chemical Theory and Com-
 840 putation* 12 (2016) 1845–1852. URL: [http://dx.doi.org/10.1021/](http://dx.doi.org/10.1021/acs.jctc.6b00049)
 841 [acs.jctc.6b00049](http://dx.doi.org/10.1021/acs.jctc.6b00049). doi:10.1021/acs.jctc.6b00049.
- 842 23. M. Khoshlessan, I. Paraskevagos, S. Jha, O. Beckstein, Parallel analysis in
 843 MDAnalysis using the Dask parallel computing library, in: Katy Huff, David
 844 Lippa, Dillon Niederhut, M. Pacer (Eds.), *Proceedings of the 16th Python in
 845 Science Conference, SciPy, Austin, TX, 2017*, pp. 64–72. doi:10.25080/
 846 shinma-7f4c6e7-00a.
- 847 24. I. Paraskevagos, A. Luckow, M. Khoshlessan, G. Chantzialexiou, T. E.
 848 Cheatham, O. Beckstein, G. Fox, S. Jha, Task-parallel analysis of molecular
 849 dynamics trajectories, in: *ICPP 2018: 47th International Conference on Parallel
 850 Processing, August 13–16, 2018, Eugene, OR, USA, Association for Comput-
 851 ing Machinery, ACM, New York, NY, USA, 2018*, p. Article No. 49.
- 852 25. P. Liu, D. K. Agrafiotis, D. L. Theobald, Fast determination of the optimal
 853 rotational matrix for macromolecular superpositions, *J Comput Chem* 31 (2010)
 854 1561–3. doi:10.1002/jcc.21439.
- 855 26. A. R. Leach, *Molecular Modelling. Principles and Applications*, Longman,
 856 1996.
- 857 27. M. Rocklin, Dask: Parallel computation with blocked algorithms and task
 858 scheduling, in: *Proceedings of the 14th Python in Science Conference, 2015*,
 859 pp. 130–136. URL: <https://github.com/dask/dask>.
- 860 28. L. D. Dalcín, R. R. Paz, P. A. Kler, A. Cosimo, Parallel distributed com-
 861 puting using python, *Advances in Water Resources* 34 (2011) 1124 – 1139.
 862 doi:10.1016/j.advwatres.2011.04.013, new Computational Meth-
 863 ods and Software Tools.

- 864 29. L. Dalcín, R. Paz, M. Storti, MPI for python, *Journal of Parallel and Distributed*
865 *Computing* 65 (2005) 1108 – 1115. doi:10.1016/j.jpdc.2005.03.010.
- 866 30. P. Garraghan, X. Ouyang, R. Yang, D. McKee, J. Xu, Straggler root-cause and
867 impact analysis for massive-scale virtualized cloud datacenters, *IEEE Transac-*
868 *tions on Services Computing* 12 (2016) 91–104. doi:10.1109/TSC.2016.
869 2611578.
- 870 31. J. Towns, T. Cockerill, M. Dahan, I. Foster, K. Gaither, A. Grimshaw, V. Hazle-
871 wood, S. Lathrop, D. Lifka, G. D. Peterson, R. Roskies, J. R. Scott, N. Wilkins-
872 Diehr, XSEDE: Accelerating scientific discovery, *Computing in Science &*
873 *Engineering* 16 (2014) 62–74. URL: doi.ieeecomputersociety.org/
874 10.1109/MCSE.2014.80. doi:10.1109/MCSE.2014.80.
- 875 32. J. A. DAILY, GAIN: DISTRIBUTED ARRAY COMPUTATION WITH
876 PYTHON, Master’s thesis, School of Electrical Engineering and Computer Sci-
877 ence, WASHINGTON STATE UNIVERSITY, 2009.
- 878 33. J. Nieplocha, B. Palmer, V. Tipparaju, M. Krishnan, H. Trease, E. Aprà, Ad-
879 vances, applications and performance of the global arrays shared memory pro-
880 gramming toolkit, *The International Journal of High Performance Computing*
881 *Applications* 20 (2006) 203–231.
- 882 34. J. Dean, S. Ghemawat, Mapreduce: Simplified data processing on large clusters,
883 in: *OSDI’04 Sixth Symposium on Operating System Design and Implementa-*
884 *tion*, 2004, pp. pp. 137–150.
- 885 35. G. Ananthanarayanan, S. Kandula, A. Greenberg, I. Stoica, Y. Lu, B. Saha,
886 E. Harris, Reining in the outliers in map-reduce clusters using mantri, in:
887 *Proceedings of the 9th USENIX Conference on Operating Systems Design and*
888 *Implementation*, OSDI’10, USENIX Association, Berkeley, CA, USA, 2010,
889 pp. 265–278.
- 890 36. J. Kyong, J. Jeon, S.-S. Lim, Improving scalability of apache spark-based
891 scale-up server through docker container-based partitioning, in: *Proceedings*
892 *of the 6th International Conference on Software and Computer Applications -*
893 *ICSCA ’17*, ACM Press, New York, USA, 2017, pp. 176–180. URL: <http://dl.acm.org/citation.cfm?doid=3056662.3056686>. doi:10.
894

- 1145/3056662.3056686.
37. K. Ousterhout, Architecting for Performance Clarity in Data Analytics Frameworks, Ph.D. thesis, EECS Department, University of California, Berkeley, Berkeley, CA, 2017. URL: <https://www2.eecs.berkeley.edu/Pubs/TechRpts/2017/EECS-2017-158.html>.
arXiv:<https://www2.eecs.berkeley.edu/Pubs/TechRpts/2017/EECS-2017-158.pdf>.
 38. A. Gittens, A. Devarakonda, E. Racah, M. Ringenburt, L. Gerhardt, J. Kottalam, J. Liu, K. Maschhoff, S. Canon, J. Chhugani, P. Sharma, J. Yang, J. Demmel, J. Harrell, V. Krishnamurthy, M. W. Mahoney, Prabhat, Matrix factorizations at scale: A comparison of scientific data analytics in spark and c++mpi using three case studies, in: IEEE International Conference on Big Data (Big Data), 2016, pp. 204–213. URL: <http://ieeexplore.ieee.org/document/7840606/>. doi:10.1109/BigData.2016.7840606.
 39. H. Yang, X. Liu, S. Chen, Z. Lei, H. Du, C. Zhu, Improving Spark performance with MPTE in heterogeneous environments, in: 2016 International Conference on Audio, Language and Image Processing (ICALIP), IEEE, 2016, pp. 28–33. URL: <http://ieeexplore.ieee.org/document/7846627/>. doi:10.1109/ICALIP.2016.7846627.
 40. E. Schmidt, G. DeMichillie, F. Perry, T. Akidau, D. Halperin, Large-scale data analysis at cloud scale, in: Symposium on Frontiers in Big Data, 2016.
 41. T.-D. Phan, Energy-efficient Straggler Mitigation for Big Data Applications on the Clouds, Ph.D. thesis, École normale supérieure de Renne, 2017.
 42. Q. Chen, C. Liu, Z. Xiao, Improving mapreduce performance using smart speculative execution strategy, IEEE Transactions on Computers 63 (2014) 954–967. doi:10.1109/TC.2013.15.
 43. B. Xie, J. Chase, D. Dillow, O. Drokin, S. Klasky, S. Oral, N. Podhorszki, Characterizing output bottlenecks in a supercomputer, in: Proceedings of the International Conference on High Performance Computing, Networking, Storage and Analysis, SC '12, IEEE Computer Society Press, Los Alamitos, CA, USA, 2012, pp. 8:1–8:11. URL: <http://dl.acm.org/citation.cfm?id=2388996.2389007>.

- 926 44. J. Rosen, B. Zhao, Fine-Grained Micro-Tasks for MapRe-
927 duce Skew-Handling, Technical Report, EECS, UC Berkeley,
928 2012. URL: [https://pdfs.semanticscholar.org/3617/](https://pdfs.semanticscholar.org/3617/916adb83f33f8df7d0b3bfc23d0de80da9b7.pdf)
929 [916adb83f33f8df7d0b3bfc23d0de80da9b7.pdf](https://pdfs.semanticscholar.org/3617/916adb83f33f8df7d0b3bfc23d0de80da9b7.pdf).
- 930 45. Y. Kwon, M. Balazinska, B. Howe, J. Rolia, Skewtune: Mitigating skew in
931 mapreduce applications, pages 25-36, in: SIGMOD'12, SIGMOD '12 Proceed-
932 ings of the 2012 ACM SIGMOD International Conference on Management of
933 Data, 2012, pp. Pages 25–36. doi:10.1145/2213836.2213840.
- 934 46. K. Ousterhout, R. Rasti, S. Ratnasamy, S. Shenker, B.-G. Chun, Making sense
935 of performance in data analytics frameworks, in: NSDI'15 Proceedings of the
936 12th USENIX Conference on Networked Systems Design and Implementation,
937 ISBN: 978-1-931971-218, 2015, pp. Pages 293–307.
- 938 47. B. Abdul-Wahid, H. Feng, D. Rajan, R. Costaouec, E. Darve, D. Thain, J. A.
939 Izaguirre, Awe-wq, fast-forwarding molecular dynamics using the accelerated
940 weighted ensemble, *Journal of Chemical Information and Modeling* 54 (2014)
941 3033–3043.
- 942 48. G. Wu, H. Song, D. Lin, A scalable parallel framework for microstructure
943 analysis of large-scale molecular dynamics simulations data, *Computational*
944 *Materials Science* 144 (2018) 322–330.
- 945 49. T. Tu, C. A. Rendleman, P. J. Miller, F. Sacerdoti, R. O. Dror, D. E. Shaw,
946 Accelerating parallel analysis of scientific simulation data via zazen, in: 8th
947 USENIX Conference on File and Storage Technologies, San Jose, CA, USA,
948 http://www.usenix.org/events/fast10/tech/full_papers/tu.pdf, 2010, pp. 129–142.
- 949 50. J. E. Stone, B. Isralewitz, K. Schulten, Early experiences scaling vmd molecular
950 visualization and analysis jobs on blue waters, in: Proceedings of the 2013
951 Extreme Scaling Workshop (Xsw 2013), IEEE Computer Society, Washington,
952 DC, USA, 2013, pp. 43–50.
- 953 51. A. Shkurtia, R. Goni, P. Andrio, E. Breitmoserd, I. Bethuned, M. Orozco, C. A.
954 Laughtona, pyPczip: A PCA-based toolkit for compression and analysis of
955 molecular simulation data, *SoftwareX* 5 (2016) 44–50.
- 956 52. P. Malakar, C. Knight, T. Munson, V. Vishwanath, M. E. Papka, Scalable in situ

- analysis of molecular dynamics simulations, in: ISAV'17 Proceedings of the In Situ Infrastructures on Enabling Extreme-Scale Analysis and Visualization, 2017, pp. 1–6.
53. T. Johnston, B. Zhang, A. Liwo, S. Crivelli, M. Taufer, *In situ* data analytics and indexing of protein trajectories, *J Comput Chem* 38 (2017) 1419–1430. doi:10.1002/jcc.24729.
54. B. R. Brooks, C. L. Brooks III., A. D. J. Mackerell, L. Nilsson, R. J. Petrella, B. Roux, Y. Won, G. Archontis, C. Bartels, S. Boresch, A. Caffisch, L. Caves, Q. Cui, A. R. Dinner, M. Feig, S. Fischer, J. Gao, M. Hodoscek, W. Im, K. Kuczera, T. Lazaridis, J. Ma, V. Ovchinnikov, E. Paci, R. W. Pastor, C. B. Post, J. Z. Pu, M. Schaefer, B. Tidor, R. M. Venable, H. L. Woodcock, X. Wu, W. Yang, D. M. York, M. Karplus, CHARMM: the biomolecular simulation program., *J. Comp. Chem.* 30 (2009) 1545–1614. doi:10.1002/jcc.21287.
55. M. J. Abraham, T. Murtola, R. Schulz, S. Páll, J. C. Smith, B. Hess, E. Lindahl, GROMACS: High performance molecular simulations through multi-level parallelism from laptops to supercomputers, *SoftwareX* 1–2 (2015) 19 – 25. doi:10.1016/j.softx.2015.06.001.
56. D. A. Case, T. E. Cheatham, 3rd, T. Darden, H. Gohlke, R. Luo, K. M. Merz, Jr, A. Onufriev, C. Simmerling, B. Wang, R. J. Woods, The amber biomolecular simulation programs, *J Comput Chem* 26 (2005) 1668–1688. doi:10.1002/jcc.20290.
57. J. Phillips, R. Braun, W. Wang, J. Gumbart, E. Tajkhorshid, E. Villa, C. Chipot, R. Skeel, L. Kale, K. Schulten, Scalable molecular dynamics with NAMD, *J Comput Chem* 26 (2005) 1781–1802. doi:10.1002/jcc.20289.
58. S. K. Burley, H. M. Berman, C. Bhikadiya, C. Bi, L. Chen, L. D. Costanzo, C. Christie, J. M. Duarte, S. Dutta, et al., Protein Data Bank: the single global archive for 3D macromolecular structure data, *Nucleic Acids Research* 47 (2018) D520–D528. URL: <http://dx.doi.org/10.1093/nar/gky949>. doi:10.1093/nar/gky949.
59. S. Van Der Walt, S. C. Colbert, G. Varoquaux, The numpy array: a structure for efficient numerical computation, *Computing in Science & Engineering* 13


- (2011) 22–30. doi:10.1109/MCSE.2011.37.
60. D. L. Theobald, Rapid calculation of RMSDs using a quaternion-based characteristic polynomial, *Acta Crystallogr A* 61 (2005) 478–80. doi:10.1107/S0108767305015266.
61. J. Daily, A. Vishnu, B. Palmer, H. van Dam, D. Kerbyson, On the suitability of MPI as a PGAS runtime, in: 2014 21st International Conference on High Performance Computing (HiPC), 2014, pp. 1–10. doi:10.1109/HiPC.2014.7116712.
62. J. Nieplocha, R. J. Harrison, R. J. Littlefield, Global arrays: A non-uniform-memory-access programming model for high-performance computers, *Journal of Supercomputing* 10 (1996) 169–189.
63. A. Collette, Python and hdf5, in: M. Blanchette, R. Roumeliotis (Eds.), Python and HDF5, O’Reilly Media, Inc., 1005 Gravenstein Highway North, Sebastopol, CA 95472., 2014.
64. S. L. Seyler, O. Beckstein, Sampling of large conformational transitions: Adenylate kinase as a testing ground, *Molec. Simul.* 40 (2014) 855–877. doi:10.1080/08927022.2014.919497.
65. S. Seyler, O. Beckstein, Molecular dynamics trajectory for benchmarking MDAnalysis, online, 2017. URL: https://figshare.com/articles/Molecular_dynamics_trajectory_for_benchmarking_MDAnalysis/5108170. doi:10.6084/m9.figshare.5108170.
66. E. Lindahl, B. Hess, D. van der Spoel, Gromacs 3.0: A package for molecular simulation and trajectory analysis, *J. Mol. Mod.* 7 (2001) 306–317. URL: <http://www.gromacs.org>. doi:10.1007/s008940100045.
67. D. Spångberg, D. S. D. Larsson, D. van der Spoel, Trajectory NG: portable, compressed, general molecular dynamics trajectories, *J Mol Model* 17 (2011) 2669–85. doi:10.1007/s00894-010-0948-5.
68. D. E. Shaw, R. O. Dror, J. K. Salmon, J. P. Grossman, K. M. Mackenzie, J. A. Bank, C. Young, M. M. Deneroff, B. Batson, K. J. Bowers, E. Chow, M. P. Eastwood, D. J. Ierardi, J. L. Klepeis, J. S. Kuskin, R. H. Larson, K. Lindorff-Larsen, P. Maragakis, M. A. Moraes, S. Piana, Y. Shan, B. Towles, Millisecond-

- 1019 scale molecular dynamics simulations on anton, in: SC '09: Proceedings of the
1020 Conference on High Performance Computing Networking, Storage and Anal-
1021 ysis, ACM, New York, NY, USA, 2009, pp. 1–11. doi:10.1145/1654059.
1022 1654099.
- 1023 69. D. E. Shaw, J. P. Grossman, J. A. Bank, B. Batson, J. A. Butts, J. C. Chao,
1024 M. M. Deneroff, R. O. Dror, A. Even, C. H. Fenton, A. Forte, J. Gagliardo,
1025 G. Gill, B. Greskamp, C. R. Ho, D. J. Ierardi, L. Iserovich, J. S. Kuskin,
1026 R. H. Larson, T. Layman, L. Lee, A. K. Lerer, C. Li, D. Killebrew, K. M.
1027 Mackenzie, S. Y. Mok, M. A. Moraes, R. Mueller, L. J. Nociolo, J. L. Peti-
1028 colas, T. Quan, D. Ramot, J. K. Salmon, D. P. Scarpazza, U. B. Schafer,
1029 N. Siddique, C. W. Snyder, J. Spengler, P. T. P. Tang, M. Theobald, H. Toma,
1030 B. Towles, B. Vitale, S. C. Wang, C. Young, Anton 2: Raising the bar for
1031 performance and programmability in a special-purpose molecular dynamics su-
1032 percomputer, in: SC '14: Proceedings of the International Conference for High
1033 Performance Computing, Networking, Storage and Analysis, 2014, pp. 41–53.
1034 doi:10.1109/SC.2014.9.
- 1035 70. R. Salomon-Ferrer, A. W. Götz, D. Poole, S. Le Grand, R. C. Walker, Rou-
1036 tine microsecond molecular dynamics simulations with amber on gpus. 2. ex-
1037 plicit solvent particle mesh ewald, *Journal of Chemical Theory and Computa-*
1038 *tion* 9 (2013) 3878–3888. URL: [http://pubs.acs.org/doi/abs/10.](http://pubs.acs.org/doi/abs/10.1021/ct400314y)
1039 [1021/ct400314y](http://pubs.acs.org/doi/abs/10.1021/ct400314y). doi:10.1021/ct400314y.
- 1040 71. J. Glaser, T. D. Nguyen, J. A. Anderson, P. Lui, F. Spiga, J. A. Millan, D. C.
1041 Morse, S. C. Glotzer, Strong scaling of general-purpose molecular dynamics
1042 simulations on gpus, *Computer Physics Communications* 192 (2015) 97–107.
1043 URL: [http://www.sciencedirect.com/science/article/](http://www.sciencedirect.com/science/article/pii/S0010465515000867)
1044 [pii/S0010465515000867](http://www.sciencedirect.com/science/article/pii/S0010465515000867). doi:dx.doi.org/10.1016/j.cpc.
1045 2015.02.028.
- 1046 72. A. Choudhary, W. keng Liao, K. Gao, A. Nisar, R. Ross, R. Thakur, R. Latham,
1047 Scalable i/o and analytics, *Journal of Physics: Conference Series* 180 (2009).
- 1048 73. S. W. Son, S. Sehrish, W. keng Liao, R. Oldfield, A. Choudhary, Reducing i/o
1049 variability using dynamic i/o path characterization in petascale storage systems,


- 1050 Journal of Supercomputing 73 (2017) pp 2069–2097.
- 1051 74. K.-W. Lin, J. Chou, S. Byna, K. W. and, Optimizing fast query performance on
1052 Lustre file system, in: SSDBM Proceedings of the 25th International Confer-
1053 ence on Scientific and Statistical Database Management, Article No. 29, 2013.
- 1054 75. G. Lockwood, 2017, What is so bad about POSIX I/O?, URL: [https://](https://www.nextplatform.com/2017/09/11/whats-bad-posix-io/)
1055 www.nextplatform.com/2017/09/11/whats-bad-posix-io/.
- 1056 76. J. Mache, V. Lo, S. Garg, The impact of spatial layout of jobs on
1057 I/O hotspots in mesh networks, Journal of Parallel and Distributed Com-
1058 puting 65 (2005) 1190 – 1203. URL: [http://www.sciencedirect.](http://www.sciencedirect.com/science/article/pii/S0743731505001048)
1059 [com/science/article/pii/S0743731505001048](http://www.sciencedirect.com/science/article/pii/S0743731505001048). doi:10.1016/
1060 j.jpdc.2005.04.020, design and Performance of Networks for Super-,
1061 Cluster-, and Grid-Computing Part I.
- 1062 77. K. A. Brown, N. Jain, S. Matsuoka, M. Schulz, A. Bhatele, Interference be-
1063 tween I/O and MPI traffic on fat-tree networks, in: Proceedings of the 47th
1064 International Conference on Parallel Processing, ICPP 2018, ACM, New York,
1065 NY, USA, 2018, pp. 7:1–7:10. URL: [http://doi.acm.org/10.1145/](http://doi.acm.org/10.1145/3225058.3225144)
1066 [3225058.3225144](http://doi.acm.org/10.1145/3225058.3225144). doi:10.1145/3225058.3225144.
- 1067 78. P. de Buyl, P. H. Colberg, F. Höfling, H5MD: A structured, efficient, and
1068 portable file format for molecular data, Computer Physics Communications
1069 185 (2014) 1546 – 1553. doi:10.1016/j.cpc.2014.01.018.

1070 Appendix A. Additional Data


1071 Figure A.11 shows performance of the RMSD task on *PSC Bridges*.




(a) Scaling total



(b) Speed-up



(c) Scaling for different components




(d) Time comparison on different parts of the calculations per MPI rank (example)

Figure A.11: *PSC Bridges*: Performance of the RMSD task. Results are communicated back to rank 0. Five independent repeats were performed to collect statistics. (a-c) The error bars show standard deviation with respect to the mean. In serial, there is no communication and hence no data point is shown for $N = 1$ in (c). (d) Compute t_{comp} , read I/O t_{IO} , communication t_{comm} , ending the for loop $t_{\text{end_loop}}$, opening the trajectory $t_{\text{opening_trajectory}}$, and overheads $t_{\text{overhead1}}$, $t_{\text{overhead2}}$ per MPI rank; see Table 3 for definitions. These are data from one run of the five repeats. MPI ranks 0 and 70 are stragglers.


1072 Figure A.12 shows performance of the RMSD task on *LSU SuperMIC*.

1073 Figure A.13 shows comparison of the parallel efficiency of the RMSD task between
1074 different test cases on *SDSC Comet*, *PSC Bridges*, and *LSU SuperMIC*.


1075 Figure A.14 shows how RMSD task scales with the increase in the number of cores
1076 when the trajectories are split using *Global Arrays* for communication compared to




(a) Scaling total



(b) Speed-up



(c) Scaling for different components



(d) Time comparison on different parts of the calculations per MPI rank (example)

Figure A.12: *LSU SuperMIC*: Performance of the RMSD task with MPI. Results are communicated back to rank 0. Five independent repeats were performed to collect statistics. (a-c) The error bars show standard deviation with respect to mean. In serial, there is no communication and hence the data points for $N = 1$ are not shown in (c). (d) Compute t_{comp} , read I/O t_{IO} , communication t_{comm} , ending the for loop $t_{\text{end_loop}}$, opening the trajectory $t_{\text{opening_trajectory}}$, and overheads $t_{\text{overhead1}}$, $t_{\text{overhead2}}$ per MPI rank; see Table 3 for definitions. These are data from one run of the five repeats.

1077 using MPI for communications on *LSU SuperMIC*.

1078 Appendix B. Effect of $R_{\text{comp/comm}}$ on Performance

1079 In Section 6.3, we improved scaling limitations due to read I/O by splitting the
 1080 trajectory, but scaling remained far from ideal when MPI communication was used;
 1081 the use of *Global Arrays* lead to better scaling (see Section 6.4) because the effective
 1082 communication cost was reduced. Although we were not able to identify the reason
 1083 for the better performance of *Global Arrays* (it still uses MPI as a communicator), the

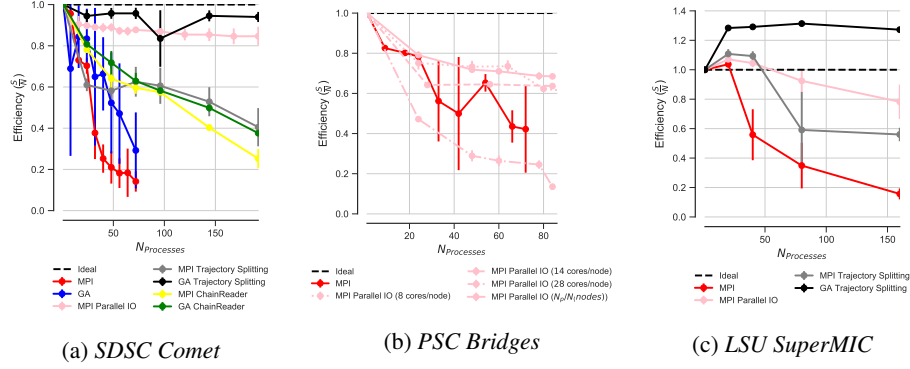


Figure A.13: Comparison of the parallel efficiency between different test cases on (a) *SDSC Comet* (data for “MPI Parallel IO” are only shown up to 192 cores for better comparison across different scenarios, see Fig. 6b for equivalent scaling data up to 384 cores), (b) *PSC Bridges*, and (c) *LSU SuperMIC*. Five repeats were performed to collect statistics and error bars show standard deviation with respect to mean.




Figure A.14: *LSU SuperMIC*: Comparison of the performance of the RMSD task with subfiling and using either MPI (“MPI”) or *Global Arrays* (“ga”) for communication. For *Global Arrays*, all ranks update the global array (`ga_put()`) and rank 0 accesses the whole RMSD array through the global memory address (`ga_get()`). Five repeats were performed to collect statistics. (a) Compute and I/O scaling versus number of processes. (b) Total time scaling versus number of processes. (c) Speed-up. (a-c) The error bars show standard deviation with respect to mean.

1084 results motivated an analysis in terms of the communication costs. In addition to the
 1085 compute to I/O ratio $R_{\text{comp}/\text{IO}}$ discussed in Section 6.2 we defined another performance
 1086 parameter called the compute to communication ratio $R_{\text{comp}/\text{comm}}$ (Eq. 8).

1087 We analyzed the data for variable workloads (see Section 6.2) in terms of the
 1088 $R_{\text{comp}/\text{comm}}$ ratio. The performance clearly depended on the $R_{\text{comp}/\text{comm}}$ ratio (Figure
 1089 B.15). Performance improved with increasing $R_{\text{comp}/\text{comm}}$ ratios (Figure B.15b and
 1090 B.15a) even if the communication time was larger (Figure B.15c). Although we still
 1091 observed stragglers due to communication at larger $R_{\text{comp}/\text{comm}}$ ratios (70× RMSD and

1092 100× RMSD), their effect on performance remained modest because the overall perfor-
 1093 mance was dominated by the compute load. Evidently, as long as overall performance
 1094 is dominated by a component such as compute that scales well, then performance prob-
 1095 lems with components such as communication will be masked and overall acceptable
 1096 performance can still be achieved (Figures B.15a and B.15b).

1097 Communication was usually not problematic within one node because of the shared
 1098 memory environment. For less than 24 processes, i.e., a single compute node on *SDSC*
 1099 *Comet*, the scaling was good and $R_{\text{comp/comm}} \gg 1$ for all RMSD loads (Figures B.15a
 1100 and B.15b). However, beyond a single compute node (> 24 cores), scaling appeared to
 1101 improve with increasing $R_{\text{comp/comm}}$ ratio while the communication overhead decreased
 1102 in importance (Figures B.15a and B.15b).

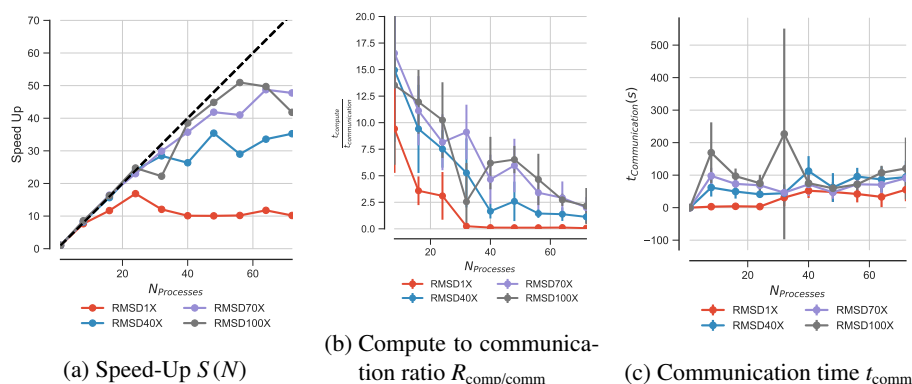


Figure B.15: Effect of the ratio of compute to communication time $R_{\text{comp/comm}}$ on scaling performance on *SDSC Comet*. (a) Scaling for different computational workloads. (Same as Figure 3a.) (b) Change in $R_{\text{comp/comm}}$ with the number of processes N for different workloads. (c) Comparison of communication time for different RMSD workloads. Five repeats were performed to collect statistics and error bars show standard deviation with respect to mean.

1103 Appendix C. Performance of the ChainReader for Split Trajectories

1104 In section 6.3.1 we showed how subfilng (splitting the trajectories) would help to
 1105 overcome I/O limitations and improve scaling. However, the number of trajectories
 1106 may not necessarily be equal to the number of processes. For example, trajectories
 1107 from MD simulations on supercomputers are often kept in small segments in individ-
 1108 ual files that need to be concatenated later to form a trajectory that can be analyzed

1109 with common tools. Such segments might be useful for subfiling but making sure that
1110 the number of processes is equal to the number of trajectory files will not always be
1111 feasible. *MDAnalysis* can transparently represent multiple trajectories as one virtual
1112 trajectory using the *ChainReader*. This feature is convenient for serial analysis when
1113 trajectories are maintained as segments. In the current implementation of *ChainReader*,
1114 each process opens all the trajectory segment files but I/O will only happen from a spe-
1115 cific block of the trajectory specific to that process only.

1116 We wanted to test if the *ChainReader* would benefit from the gains measured for
1117 the subfiling approach. Specifically, we measured if the MPI-parallelized RMSD task
1118 (with N_p ranks) would benefit if the trajectory was split into $N_{\text{seg}} = N_p$ trajectory
1119 segments, corresponding to an ideal scenario.

1120 In order to perform our experiments we had to work around an issue with the XTC
1121 format reader in *MDAnalysis* that was related to the XTC random-access functional-
1122 ity that the `MDAnalysis.coordinates.XTC.XTCReader` class provides: The
1123 Gromacs XTC format [66, 67] is a lossy-compression, XDR-based file format that was
1124 never designed for random access and the compressed format itself does not support
1125 fast random seeking. The *XTCReader* stores persistent offsets for trajectory frames
1126 to disk [18] in order to enable efficient access to random frames. These offsets will be
1127 generated automatically the first time a trajectory is opened and the offsets are stored
1128 in hidden `*.xtc_offsets.npz` files. The advantage of these persistent offset files
1129 is that after opening the trajectory for the first time, opening the same file will be very
1130 fast, and random access is immediately available. However, stored offsets can get out
1131 of sync with the trajectory they refer to. To prevent the use of stale offset data, tra-
1132 jectory file data (number of atoms, size of the file and last modification time) are also
1133 stored for validation. If any of these parameters change the offsets are recalculated. If
1134 the XTC changes but the offset file is not updated then the offset file can be detected as
1135 invalid. With *ChainReader* in parallel, each process opens all the trajectories because
1136 each process builds its own `MDAnalysis.Universe` data structure. If an invalid
1137 offset file is detected (perhaps because of wrong file modification timestamps across
1138 nodes), several processes might want to recalculate these parameters and rebuild the
1139 offset file, which can lead to a race condition. In order to avoid the race condition,

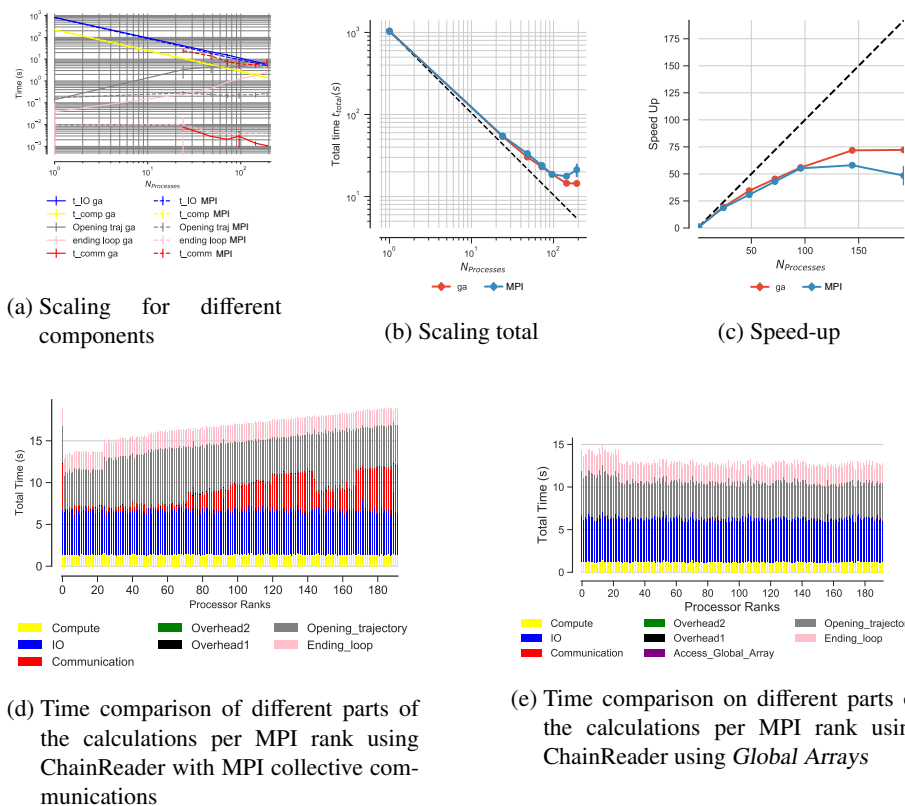


Figure C.16: Comparison on the performance of the MDAnalysis ChainReader for the RMSD task on *SDSC Comet* when the trajectories are split; for the communication step either collective MPI (“MPI”) or *Global Arrays* (“ga”) was used. In case of *Global Arrays*, all ranks update the global array (`ga_put()`) and rank 0 accesses the whole RMSD array through the global memory address (`ga_get()`). Five repeats were performed to collect statistics. (a) Compute and I/O scaling versus number of processes. (b) Total time scaling versus number of processes. (c) Speed-up. (a-c) The error bars show standard deviation with respect to the mean. (d-e) Compute t_{comp} , read I/O t_{IO} , communication t_{comm} , access to the whole global array by rank 0 $t_{\text{Access_Global_Array}}$, ending the for loop $t_{\text{end_loop}}$, opening the trajectory $t_{\text{opening_trajectory}}$, and overheads $t_{\text{overhead1}}$, $t_{\text{overhead2}}$ per MPI rank. (See Table 3 for the definitions.)

1140 we removed this check from MDAnalysis for the purpose of the measurements pre-
 1141 sented here, but this comes at the price of not checking the validity of the offset files
 1142 at all; future versions of MDAnalysis may lift this limitation. We obtained the results
 1143 for the ChainReader with this modified version of *MDAnalysis* that eliminates the race
 1144 condition by assuming that XTC index files are always valid.

1145 Figure C.16 shows the results for performance of the ChainReader for the RMSD
 1146 task using either collective MPI or Global Arrays (GA) for the communication step.
 1147 With GA good strong scaling performance was observable up to 144 cores (Figure

1148 C.16c); without GA, strong scaling plateaued for more than 92 cores. In both cases,
 1149 strong scaling performance was worse than what was achieved when each MPI process
 1150 was assigned its own trajectory segment as described in Section 6.3.1. The strong scal-
 1151 ing performance did not suffer because of the computation and the read I/O because
 1152 both t_{comp} and $t_{\text{I/O}}$ showed excellent strong scaling up to 196 cores (Figure C.16a). In-
 1153 stead the time for ending the `FOR` loop $t_{\text{end_loop}}$, which includes the time for closing the
 1154 trajectory file, and opening the trajectory $t_{\text{opening_trajectory}}$ appeared to be the scaling bot-
 1155 tleneck. These results differed from the subfiling results (section 6.3.1) where neither
 1156 $t_{\text{end_loop}}$ nor $t_{\text{opening_trajectory}}$ limited scaling (Figures 5d and 5e).

1157 Although we did not further investigate the deeper cause for the reduced scaling
 1158 performance of the ChainReader, we speculate that the primary problem is related to
 1159 each MPI rank having to open all trajectory files in their ChainReader instance even
 1160 though they will only read from a small subset. For N_p ranks and N_{seg} file segments,
 1161 in total, $N_p N_{\text{seg}}$ file opening/closing operations have to be performed. Each server that
 1162 is part of a Lustre file system can only handle a limited number of I/O requests (read,
 1163 write, stat, open, close, etc.) per second. A large number of such requests, from one
 1164 or more users and one or more jobs, can lead to contention for storage resources. For
 1165 $N_p = N_{\text{seg}} = 100$, the Lustre file system has to perform 10,000 of these operations
 1166 almost simultaneously, which might degrade performance.